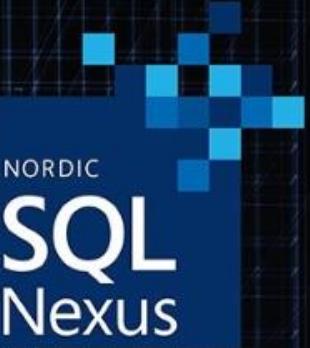


SQL Azure Database V12 - Lessons learned from the trenches



José Manuel Jurado Diaz

Azure SQL Database Support Engineer / Subject Master Expert





About me..

Jose Manuel Jurado Diaz

Azure SQL DB Support Engineer

Subject Master Expert

Speaker SQL Saturday, channel9 and
TechReady



Email: jmjurado@microsoft.com





Session Objectives And Takeaways

❑ Session Objectives:

- ❑ Raise awareness of Azure SQL DB common issues, help to improve customer's business continuity and prevent Azure SQL DB performance issues.

❑ Key Takeaway 1.

- ❑ New Options in Azure SQL DB v12

❑ Key Takeaway 2.

- ❑ New diagnostic tools for performance.

❑ Key Takeaway 3.

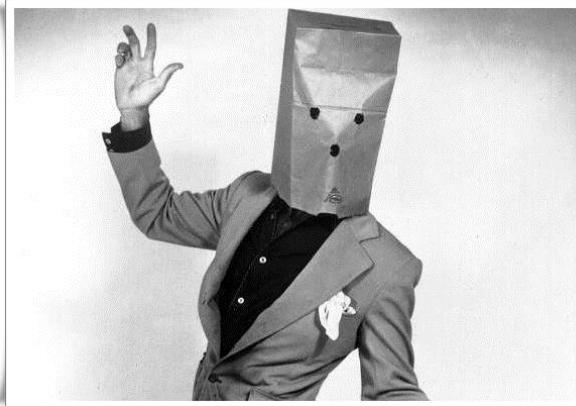
- ❑ New database models and business continuity.



Before Getting Started - Questions



Did you work with SQL Azure previously? V11 or V12?



If not, do you know what is it?



If you know, did you have any trouble or issue working with SQL Azure previously?





SQL Azure Database V12

Introduction To Azure SQL DB





Azure SQL Database – What is it?

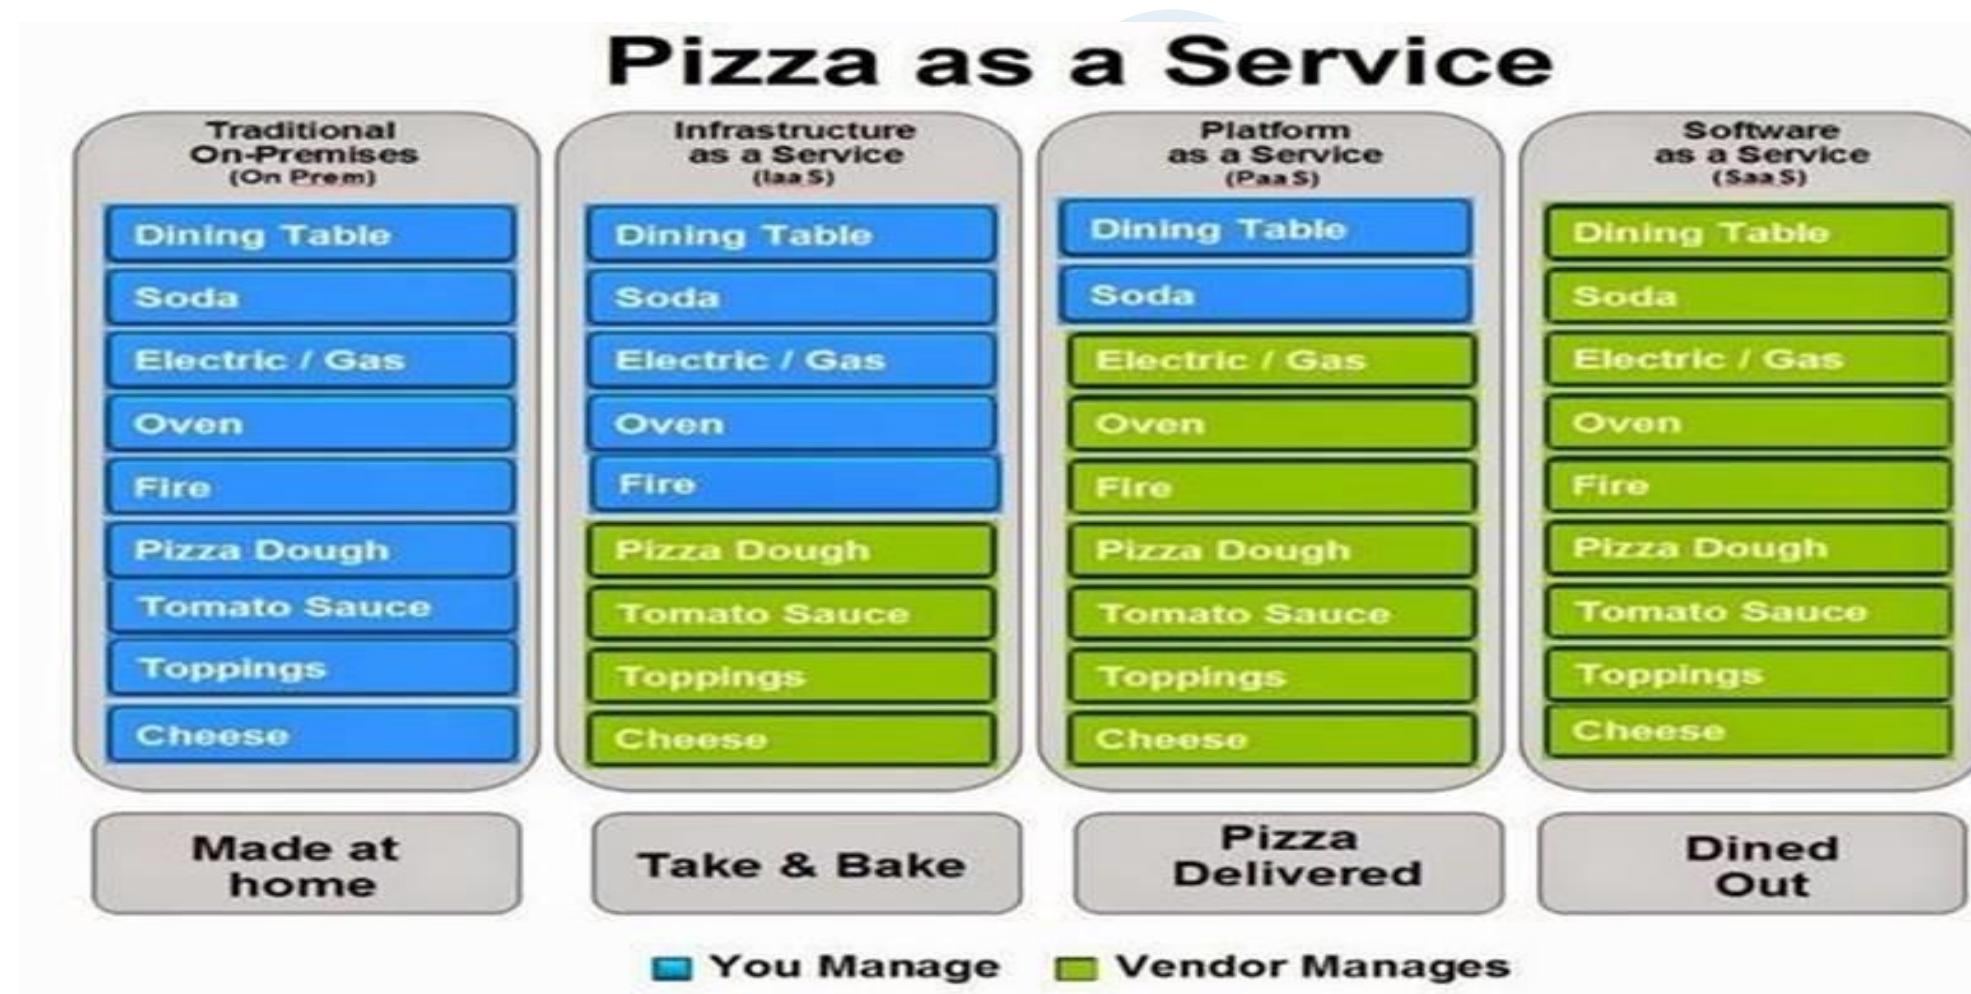


- SQL Server as Service
- Service managed by Microsoft
- Enterprise Edition options, limited only by hardware characteristics depending on the model.
- Almost 0 maintenance
- Resources oriented to use your database



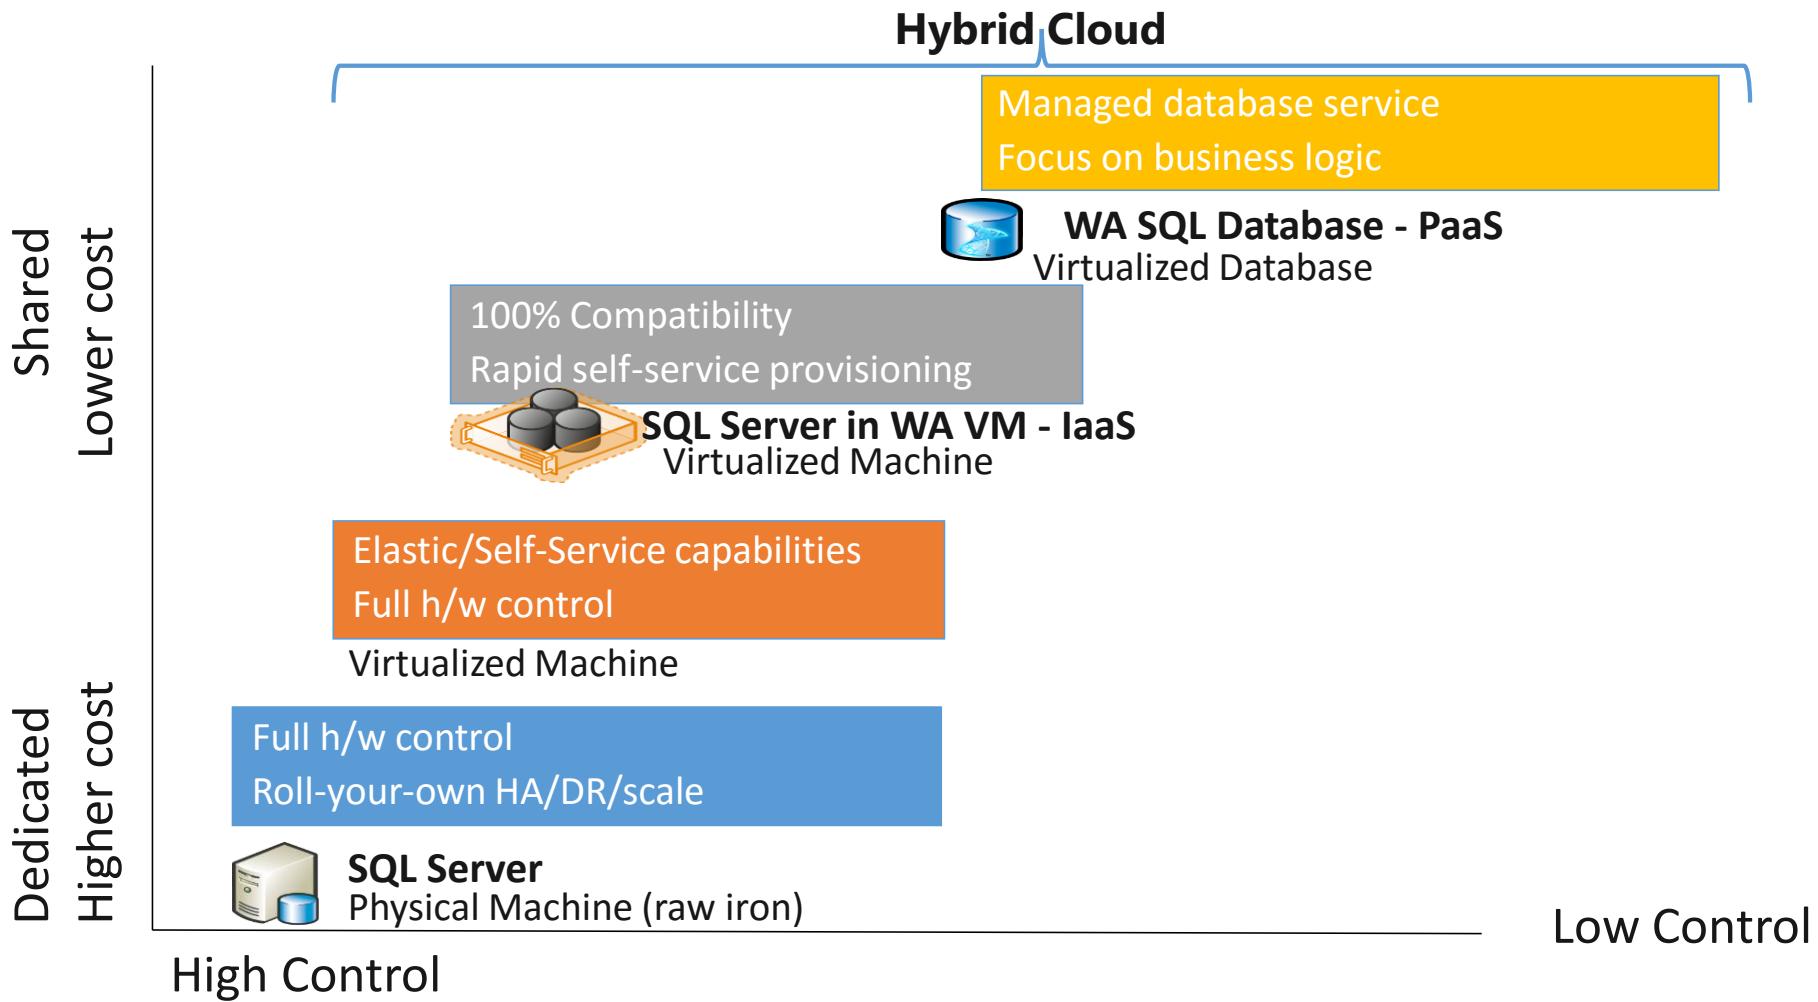


Azure SQL Database – Cloud Models





Azure SQL Database – Database Platform





Azure SQL Database – Five Goals

Simple

Easy to create.

Dashboard metrics DB.

No need to change the application.

Multiple migration tools.

Scalability

Predictive performance.

Multiple database model depending on use.

Parallel Data-Warehouse

Trusted

Azure AAD or SQL Login

Multiple security mechanisms.

Compliance

Auditing

Business Continuity

Disaster Recovery

Geo-Replication
Geo-Restore

Restore Point-in-Time.

Database Replicas.

Backups

Always Updated

Latest options available at company level.

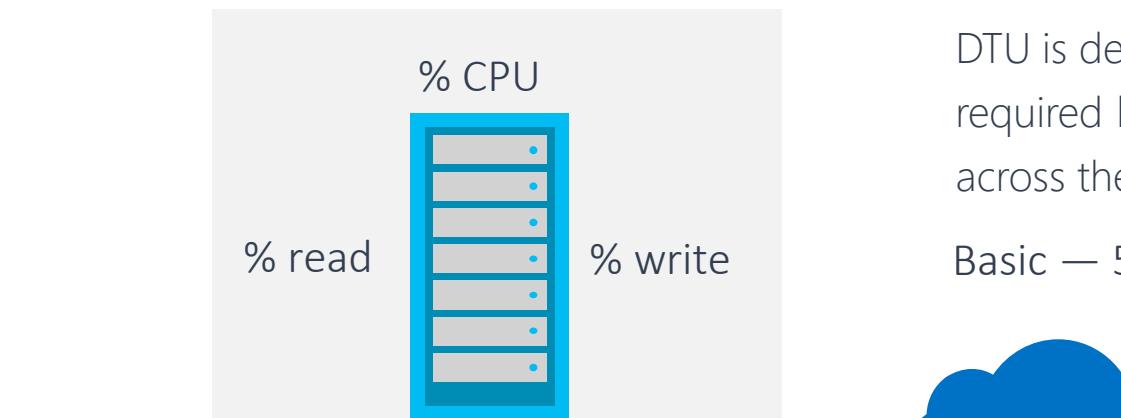
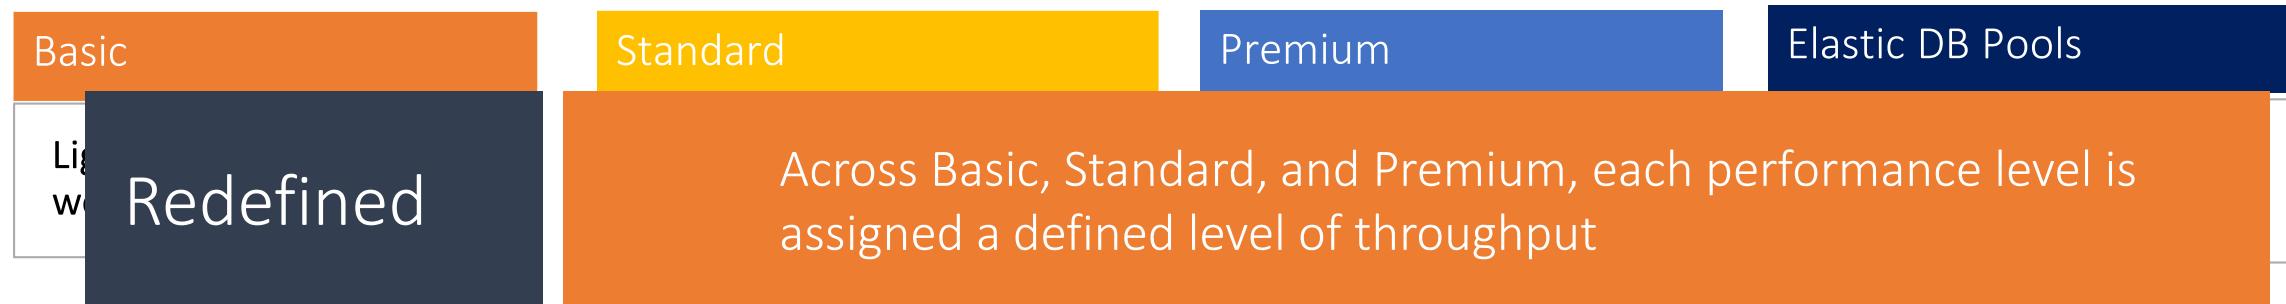
Multiple monitoring tools.



Familiarization.



Azure SQL Database – Database Level/DTUs



DTU is defined by the bounding box for the resources required by a database workload and measures power across the six performance levels.

Basic — 5 DTU	S0 — 10 DTU	P1 — 125 DTU
	S1 — 20 DTU	P2 — 250 DTU
	S2 — 50 DTU	P4 — 500 DTU

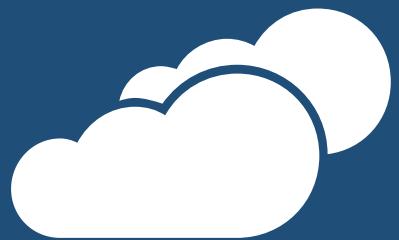
Measure of power

Introducing the Database Throughput Unit (DTU) which represents database power and replaces hardware specs

Microsoft



Azure SQL Database – Create First DB



Demo: Creating my first database





Agenda

❑ Differences with V11.

- The good and the bad

❑ Changes and new improvements in Azure SQL Database V12

- Engine
- Security

❑ High Availability and Disaster Recovery

- Geo-Replication and Geo-Restore

❑ The new database model and Elastic Tools

- Elastic Database Pools
- Elastic Query, Jobs, Scale and Transaction

❑ New Tools and connectivity improvements

- Connectivity
- Query Data Store, Query Performance Insights, Index Advisor, Extended Events, Alerts, DMVs.

❑ Q&A





SQL Azure Database V12

Differences with Azure SQL

DB v11



Differences with V11 – The Good

- ✓ No noisy neighbors.
- ✓ Instance per database non-shared.
- ✓ 20%-30% improvement.
- ✓ Direct connection and differences depending on data provider.
- ✓ Predictive performance and DTUs
- ✓ Our customer is able to restore backups.
- ✓ Dedicated Administrative Connection.
- ✓ A lot of new DMVs.
- ✓ Deleted server.
 - ✓ **Tip:** If you drop a server you wouldn't be able to reuse the name after
 - ✓ 5 days.





Differences with V11 – The Bad

- ✓ Compare Basic/Standard vs Web/Business.
 - ✓ Tip: Update stats, rebuild index, missing index. <http://blogs.msdn.com/b/sqlblog/archive/2013/11/01/do-i-need-to-upgrade-my-dba-skills-for-the-cloud.aspx>
 - ✓ Postpone the migration from V11 to V12:
<http://blogs.msdn.com/b/azuresqldbsupport/archive/2015/06/05/stopping-or-postponing-an-upgrade-to-sql-database-v12.aspx>
 - ✓ Business and Web will be migrated in automatic way started at: 12th September -
<https://azure.microsoft.com/en-us/blog/azure-sql-database-web-and-business-edition-retirement-september-12th-2015/>
- ✓ Retry-Logic is more needed than before.
- ✓ Once moved V12 not possible to go back.
- ✓ BulkCopy works but not loading entire files you could use bcp instead of.
- ✓ No more federations, now, Sharding framework. <https://azure.microsoft.com/en-us/documentation/articles/sql-database-elastic-scale-use-entity-framework-applications-visual-studio/>
- ✓ Lost HTML to manage the database.



SQL Azure Database V12

*Changes and Improvements
of Azure SQL DB*



What is new using V12 Azure SQL Engine

Engine

[Partitioning](#)

[Partial Full-Text Search](#)

Parallel Execution for >P2

[Contained database](#)

[Tables In-Memory](#)

[Temporal Tables \(Versioning\)](#)

DBCC commands

CLR Integration – Stopped

Table as Heap

[Data Compression](#)

Enhanced Session Context

Setting changes Database

[Change tracking](#)

[New functions and TSQL](#)

[Strech Database](#)

Compatibility Level to 130

[SQL Server 2016 \(Preview \)](#)

Azure SQL DB V12

Alter Database SCOPE Configuration

Available for:

[SQL Server 2016 \(Preview \)](#)

[Azure SQL DB V12](#)

Options:

Procedure cache cleaning

Maxdop

Legacy Cardinality Estimation

Parameter sniffing

Indexes

[Online Index.](#)

[Column Store Index.](#)

[XML Index.](#)

Replication

[Transactional](#)

[SQL Data-Sync](#)

New Product

[Azure SQL DataWarehouse – Parallel Data Warehouse](#)

High Availability

[Geo-Replication Offline/Online](#)





What is new using V12 Azure SQL Engine

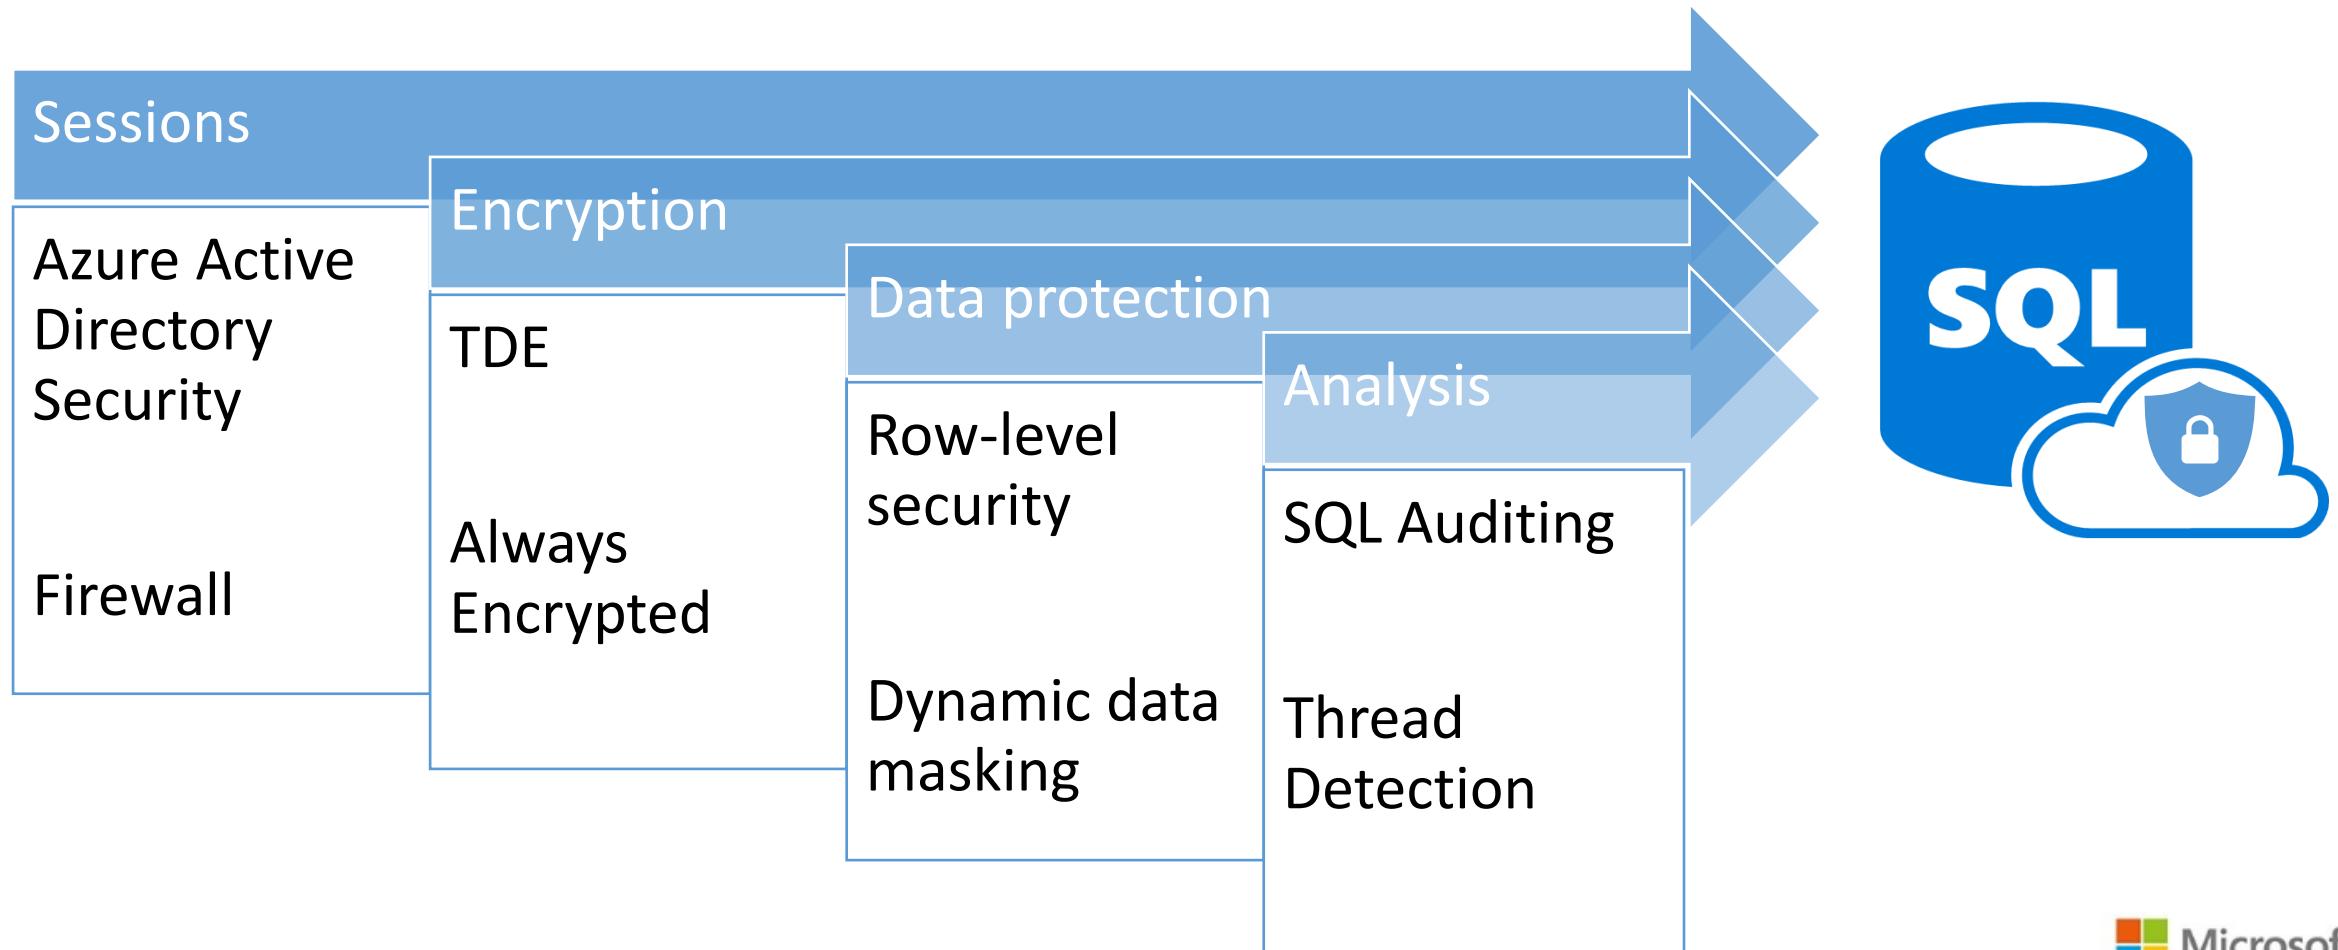


*Demo: Alter Database Scope
Configuration*





What is new using V12 Azure SQL Security





Transparent Data Encryption

Data is encrypted at rest, in flight, and while in use

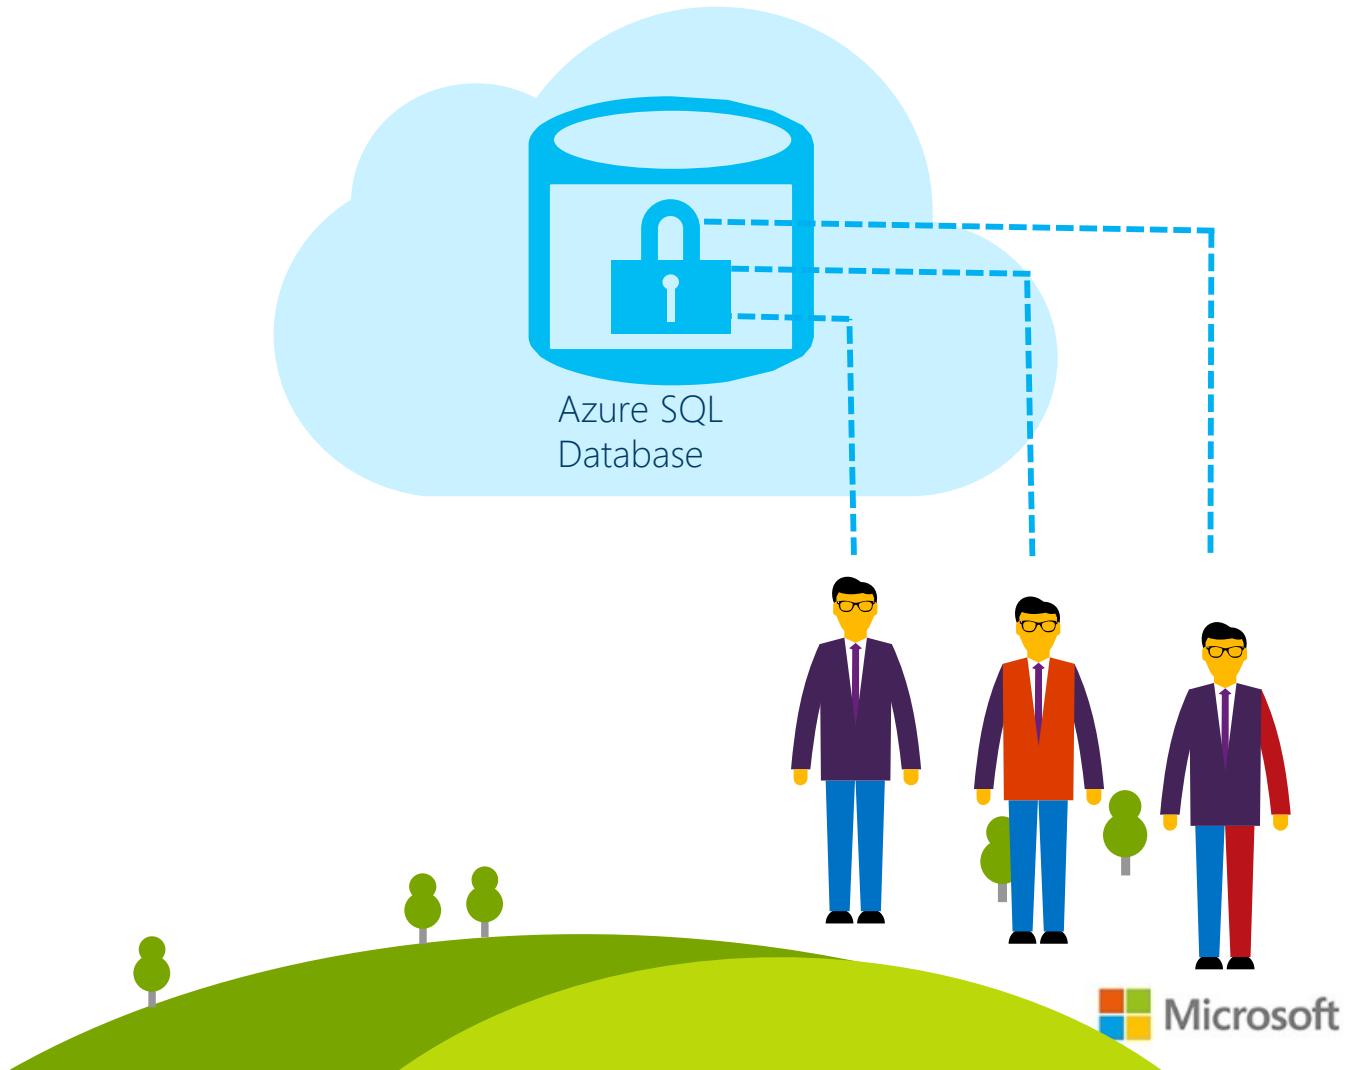
Azure SQL Database service manages your keys

You can keep application changes to a minimum

Encryption and decryption of data is done
transparently in a transparent data encryption
(TDE)-enabled client driver

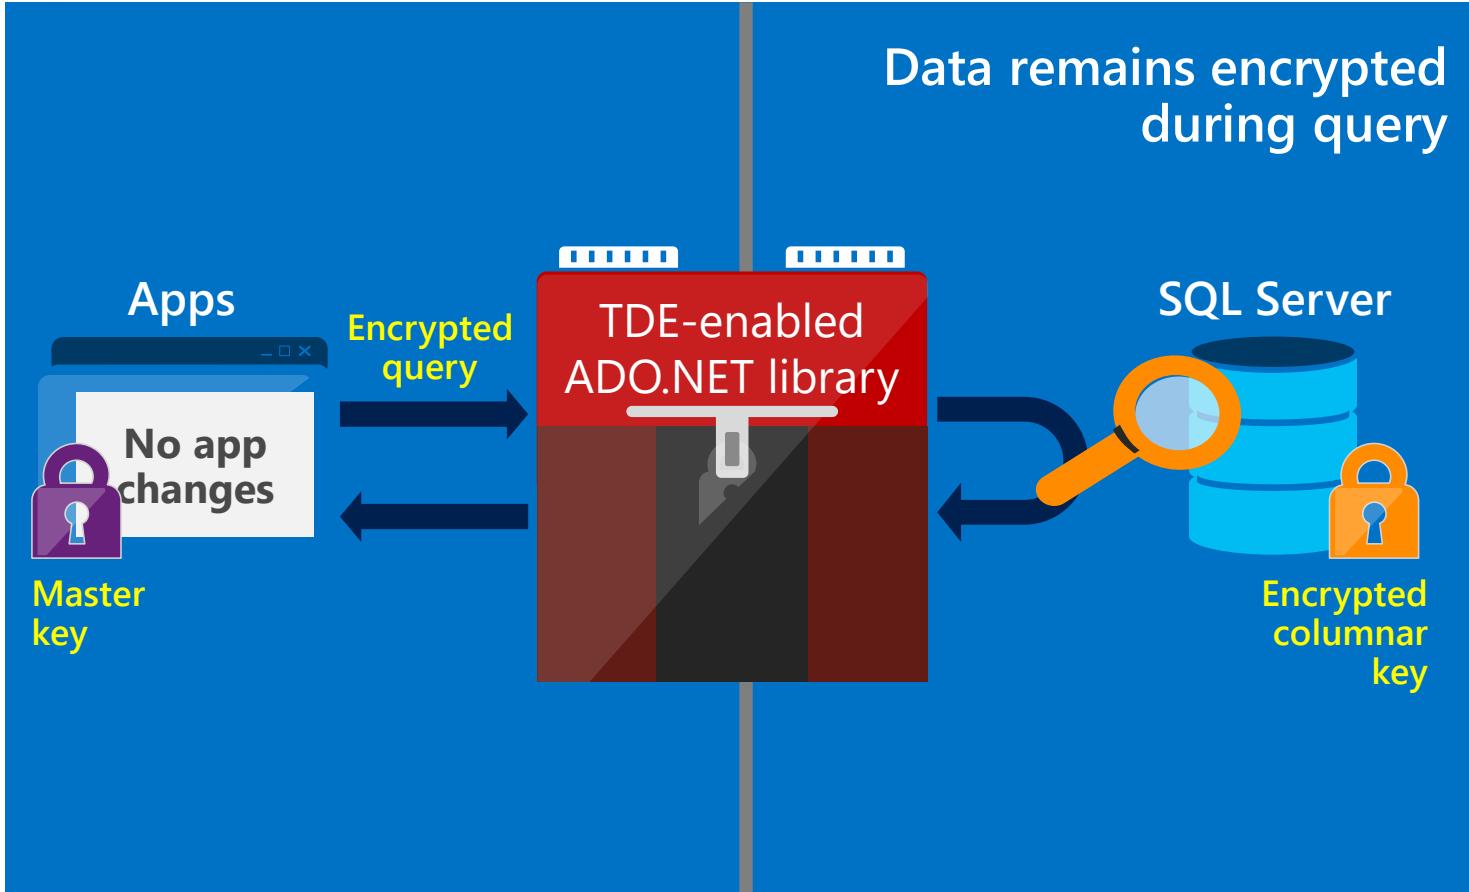
Support for equality operations (including joins) on
encrypted data

Azure manages encryption keys





Always Encrypted



Capability

Transparent client-side encryption, while SQL Server executes T-SQL queries on encrypted data

Benefits

- Sensitive data remains encrypted and queryable at all times
- Unauthorized users never have access to data or keys
- No changes to applications are necessary



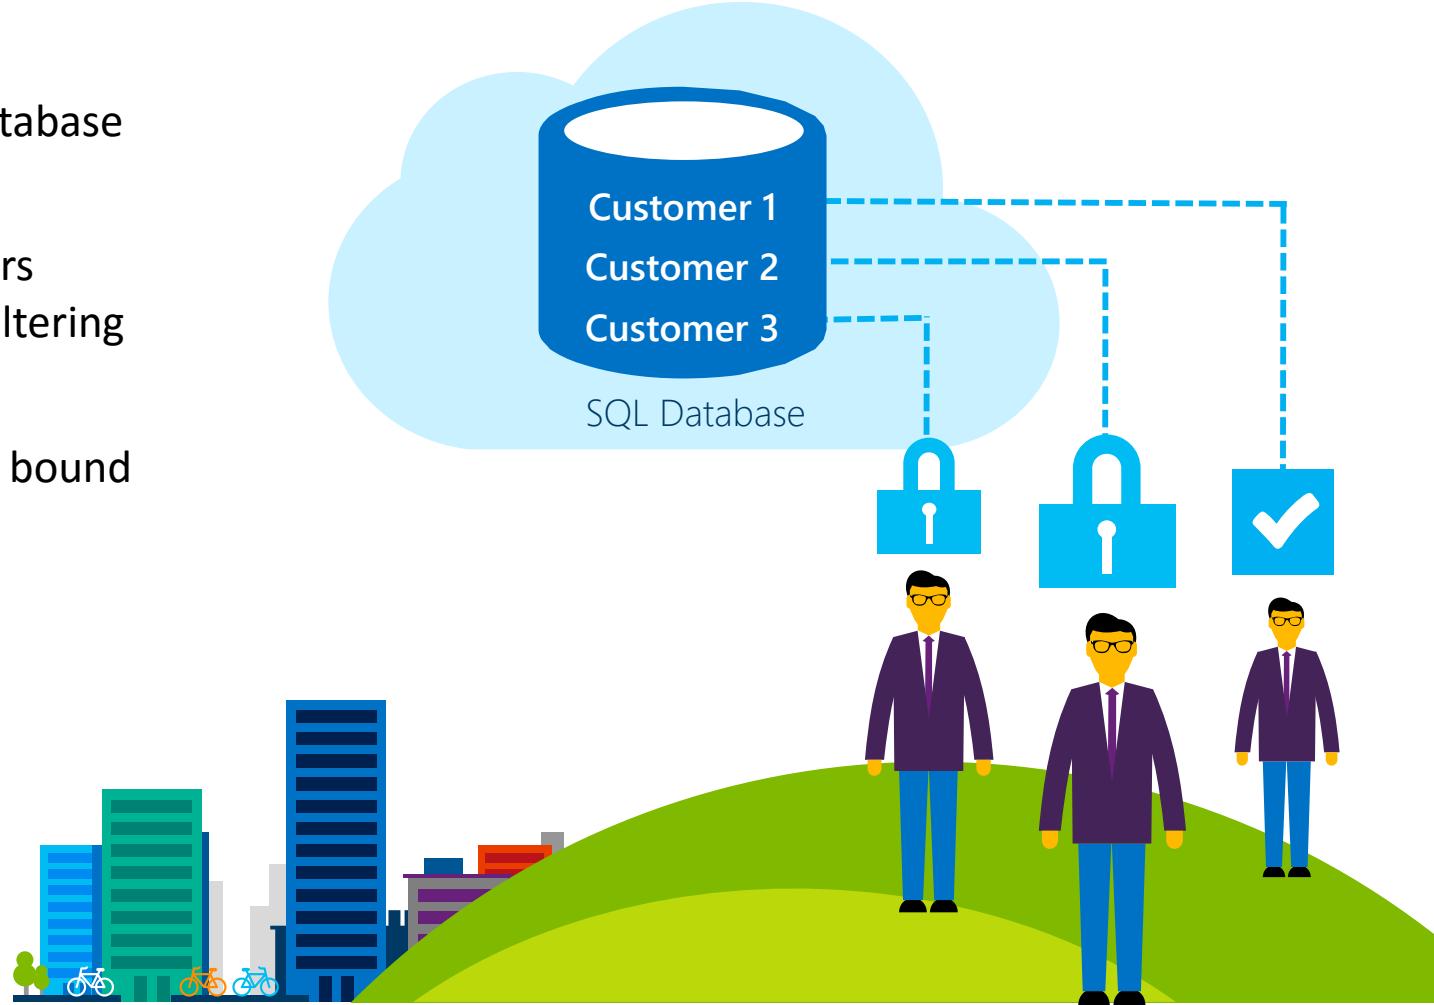


Row Level Security

Fine-grained access control over specific rows in a database table

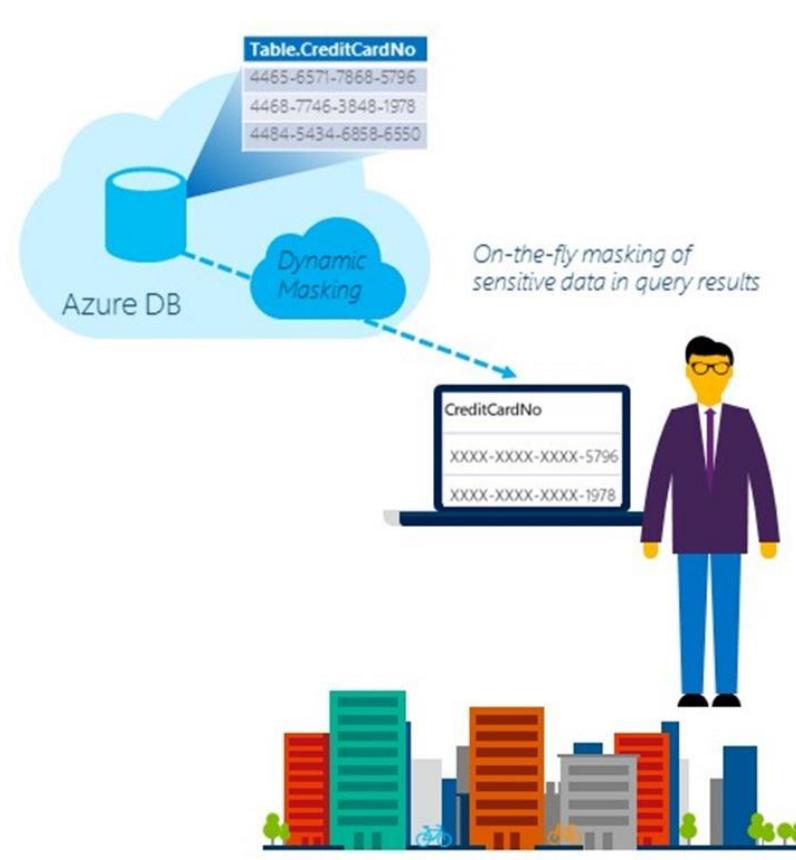
Help prevent unauthorized access when multiple users share the same tables, or to implement connection filtering in multitenant applications

Enforcement logic inside the database and schema is bound to the table





Dynamic Data Masking





Dynamic Data Masking - Example

// Data Masking

```
CREATE TABLE Contacto
(
    ID int IDENTITY PRIMARY KEY,
    Nombre varchar(100) MASKED WITH (FUNCTION = 'partial(1,"XXXXXXX",0)') NULL,
    Apellido varchar(100) NOT NULL,
    NrTlf varchar(12) MASKED WITH (FUNCTION = 'default()') NULL,
    Email varchar(100) MASKED WITH (FUNCTION = 'email()') NULL);
```

// Row Level Security

```
CREATE FUNCTION SecPred(@userId int)
RETURNS TABLE
WITH SCHEMABINDING
AS
RETURN SELECT 1 as valor WHERE @userId = user_id()
CREATE SECURITY POLICY [secpol] ADD FILTER PREDICATE
[dbo].[SecPred]([UserId]) on [dbo].[Protegido]
```





SQL Auditing

- Azure SQL Database Auditing tracks database events and writes audited events to an audit log in your Azure Storage account

Audit Records (Preview)				
Event Time	Event ID	Principal Name	Event Type	Action Status
Thu, 29 Oct 2015 08:46:58...	48ce775-3f02-49a0-8b11...	testlogin	DataAccess	Failure
Thu, 29 Oct 2015 08:46:58...	04e6fb7-b6e0-4fb2-a23a...	testlogin	Login	Success
Thu, 29 Oct 2015 07:36:39...	8bae2f31-aec0-44aa-af63...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:36...	4b0f6bd9-71c9-4d25-9a44...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:34...	9a8b2c5b-ee02-4e28-8db...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:34...	78bf75f4-4f02-410f-b6ff-c...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:34...	0bb78eb5-6699-463e-9d8...	testlogin	Login	Success
Thu, 29 Oct 2015 07:36:27...	49ca2098-ab10-4999-995f...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:20...	2ad9b130-0730-4754-83d...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:19...	747847f9-0d46-4bb7-8b2...	testlogin	Login	Success

Source	Destination	Protocol
NCPP	westeurope1-a.control.database.wind...	DNS
westeurope1-a.control.data...	NCPP	TCP
NCPP	westeurope1-a.control.database.wind...	TCP
NCPP	westeurope1-a.control.database.wind...	TDS
westeurope1-a.control.data...	NCPP	TDS
NCPP	westeurope1-a.control.database.wind...	TLS
westeurope1-a.control.data...	NCPP	TLS
westeurope1-a.control.data...	NCPP	TCP
westeurope1-a.control.data...	NCPP	TCP
NCPP	westeurope1-a.control.database.wind...	TCP
NCPP	westeurope1-a.control.database.wind...	TLS
westeurope1-a.control.data...	NCPP	TLS
NCPP	westeurope1-a.control.database.wind...	SSDP
NCPP	westeurope1-a.control.database.wind...	TDS
westeurope1-a.control.data...	NCPP	TDS
NCPP	westeurope1-a.control.database.wind...	TCP
NCPP	NCPP	DNS
westeurope1-a.control.data...	NCPP	TCP
westeurope1-a.control.data...	NCPP	TCP
NCPP	westeurope1-a.control.database.wind...	TCP
NCPP	NCPP	DNS
NCPP	datasec-weu-2-a3.cloudapp.net	TCP

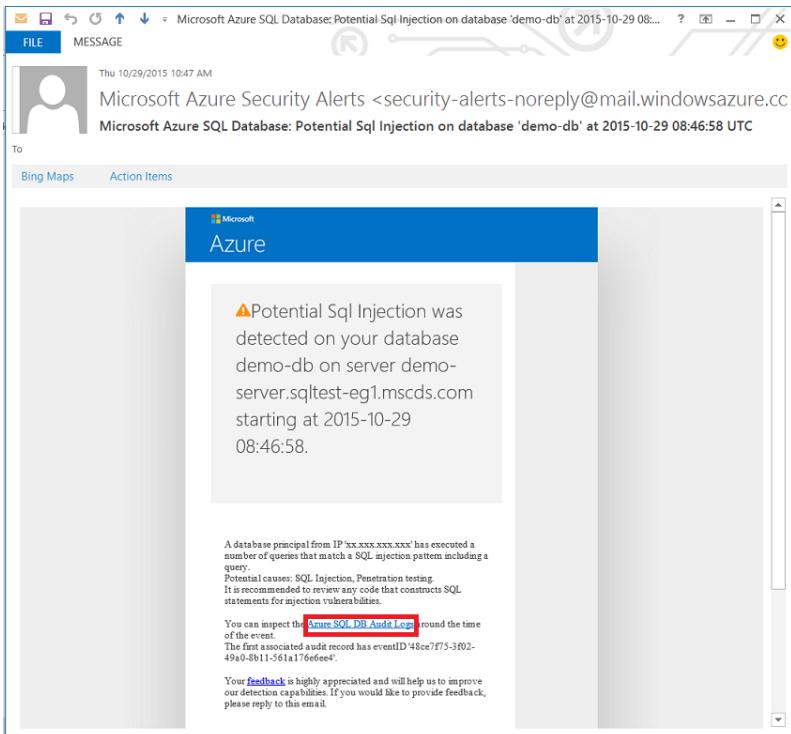
Events: Plain SQL, Parameterized SQL, Stored procedure, Logins, Transaction management





Thread Detection

- Threat Detection detects anomalous database activities indicating potential security threats to the database.

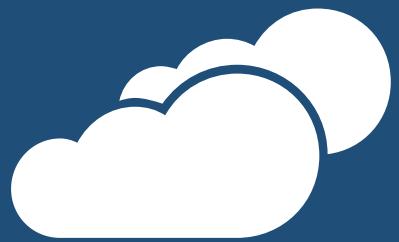


Audit Records (Preview)				
EVENT TIME	EVENT ID	PRINCIPAL NAME	EVENT TYPE	ACTION STATUS
Thu, 29 Oct 2015 08:46:58...	48ce7f75-3f02-49a0-8b11...	testlogin	DataAccess	Failure
Thu, 29 Oct 2015 08:46:58...	04e8fb7-b6e0-4fb2-a23a...	testlogin	Login	Success
Thu, 29 Oct 2015 07:36:39...	8bae2f31-aec0-44aa-af63...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:36...	4b0f6bd9-71c9-4d25-9a44...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:34...	9a8b2c5b-ee02-4e28-8db...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:34...	78bf75f4-4f02-410f-b6ff-c...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:34...	0bb78eb5-6699-463e-9d8...	testlogin	Login	Success
Thu, 29 Oct 2015 07:36:27...	49ca2098-ab10-4999-995f...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:20...	2ad9b130-0730-4754-83d...	testlogin	DataAccess	Success
Thu, 29 Oct 2015 07:36:19...	747847f9-0d46-4bb7-8b2...	testlogin	Login	Success





What is new using V12 Azure SQL Security



Demo: Row Level Security





SQL Azure Database V12

High Availability and Disaster Recovery



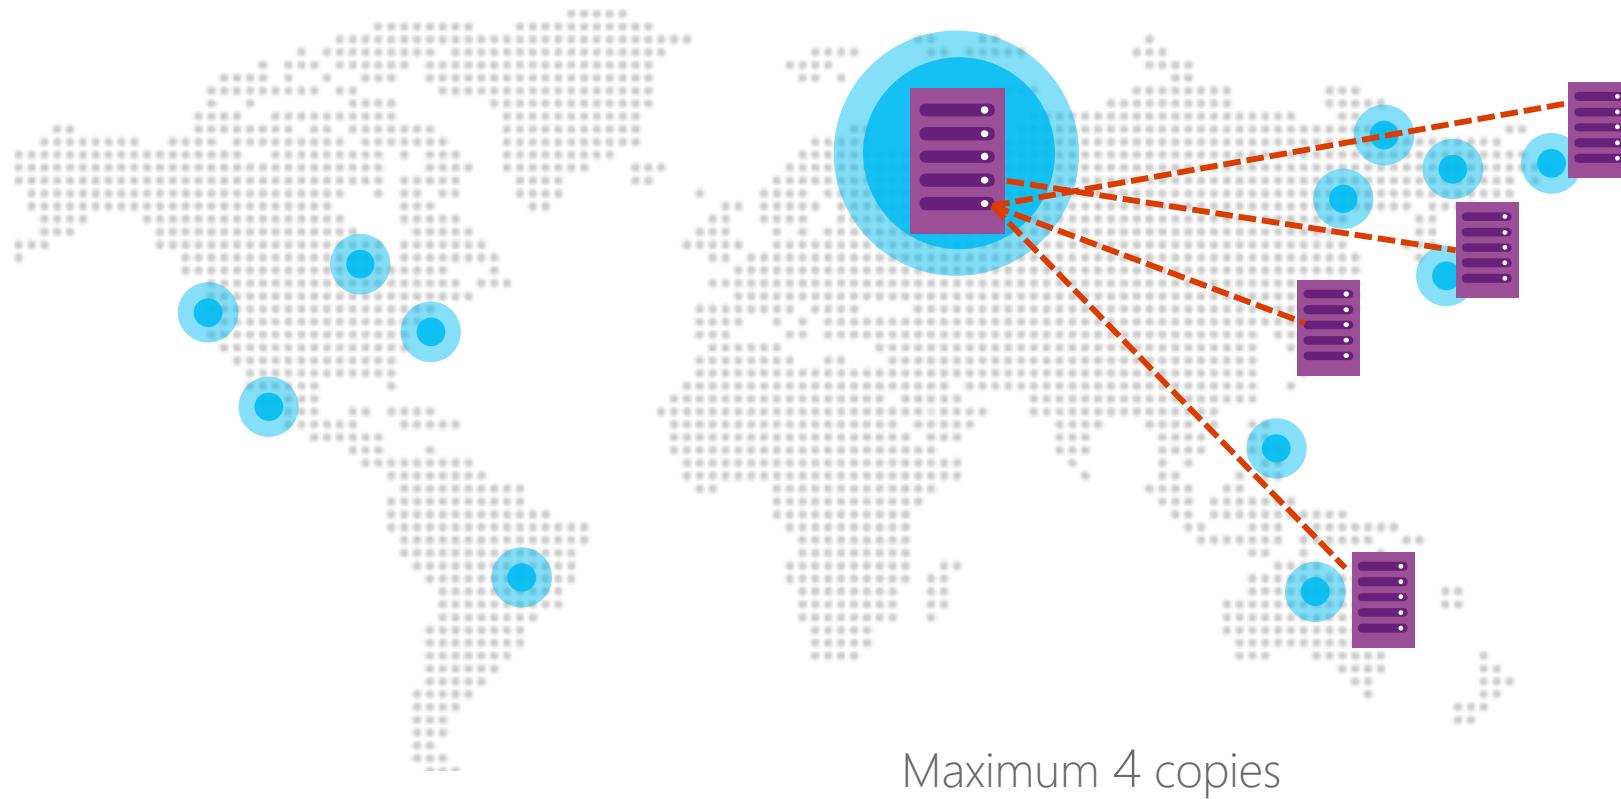


SQL Azure Database V12 – High Availability and Disaster Recovery

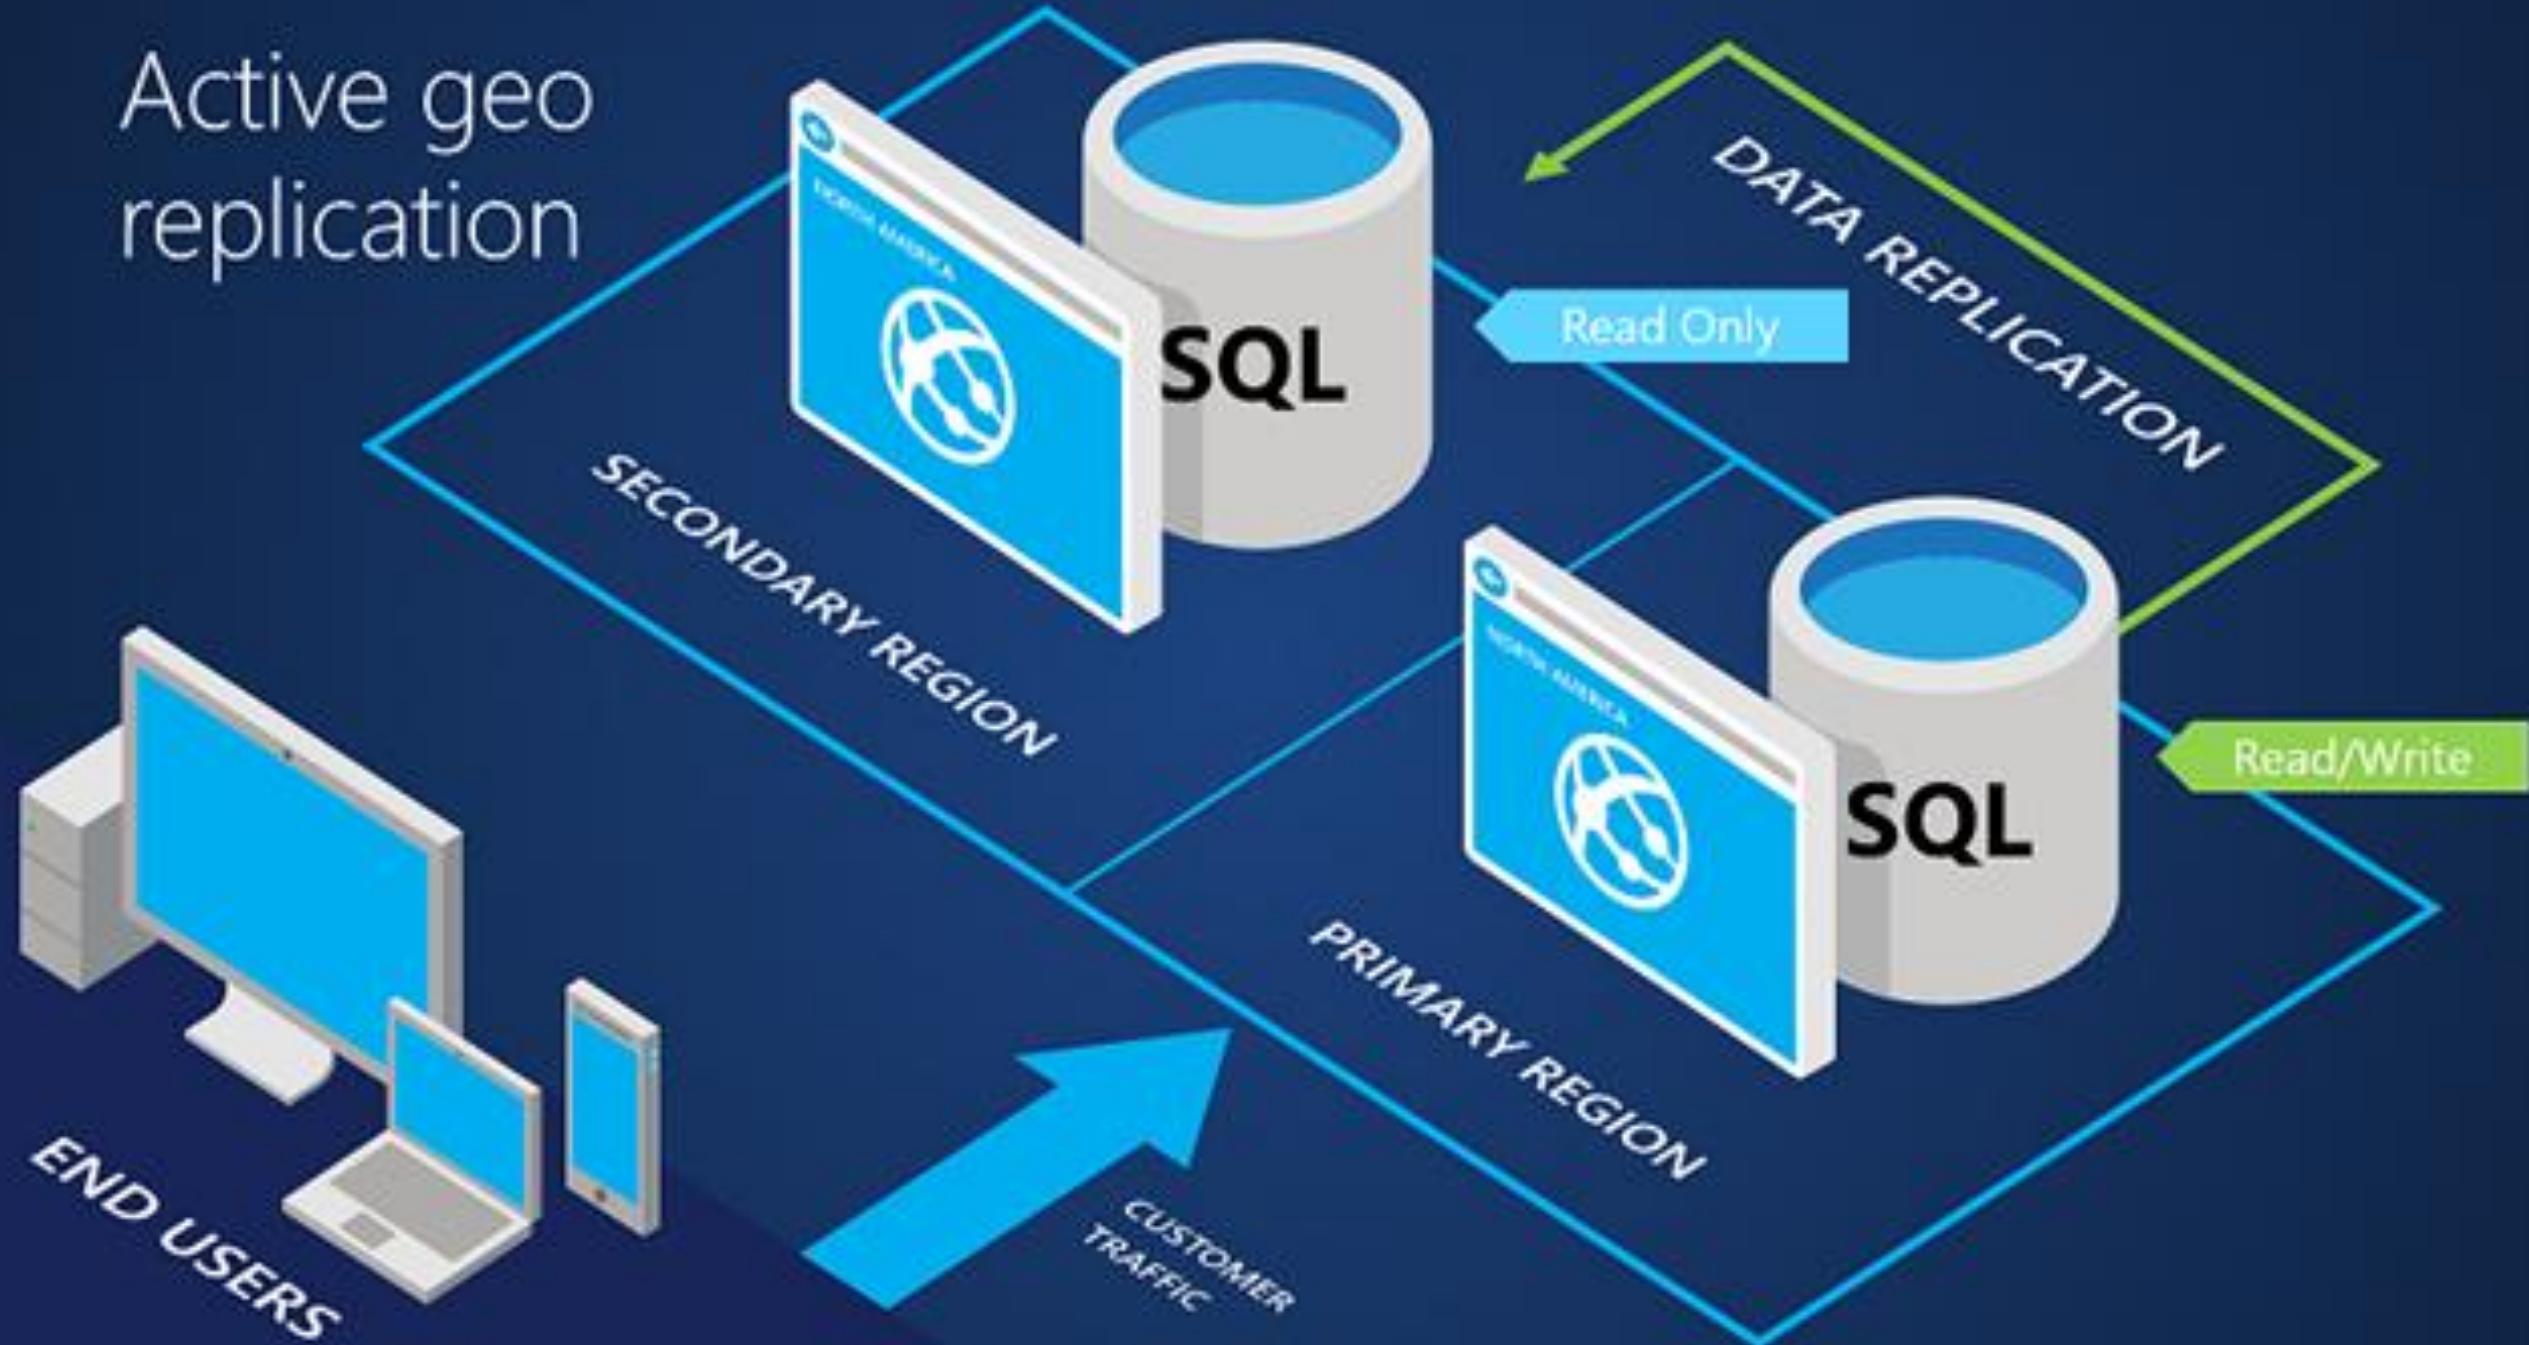
- ✓ SQL Azure HA and DR is in our DNA.
- ✓ Geo-Replication – more than 1 replica (Online).
<http://blogs.msdn.com/b/timomta/archive/2015/03/26/script-to-perform-azure-sql-premium-failover.aspx>
- ✓ Geo-Restore.
 - ✓ All database backups are geo-replicated. SQL DB supports geo-restore – eg: any database can be restored anywhere in the world. <https://azure.microsoft.com/en-us/documentation/articles/sql-database-business-continuity/>
- ✓ SQL DataSync is possible to use but not supported.
- ✓ Replication between databases. Others: [DBMoto](#)



SQL Azure Database V12 – High Availability and Disaster Recovery

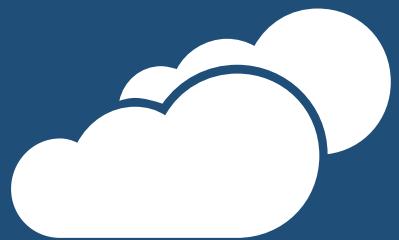


Active geo
replication





SQL Azure Database V12 – High Availability and Disaster Recovery



Demo: Implement Geo-Replication





SQL Azure Database V12

Elastic Database Pool



SQL Azure Database V12 - New Database Tiers

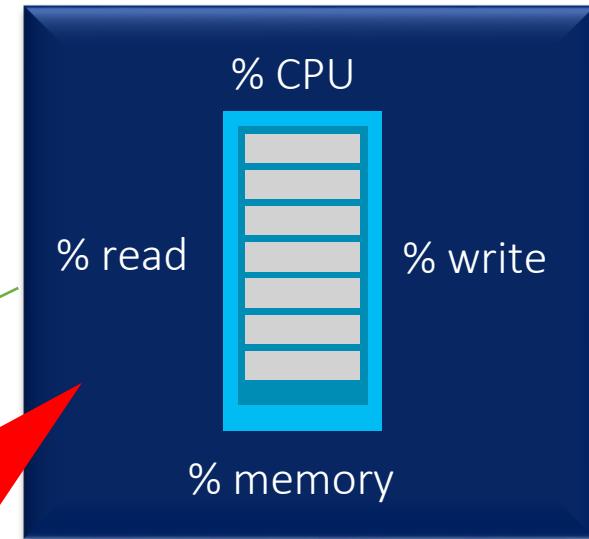
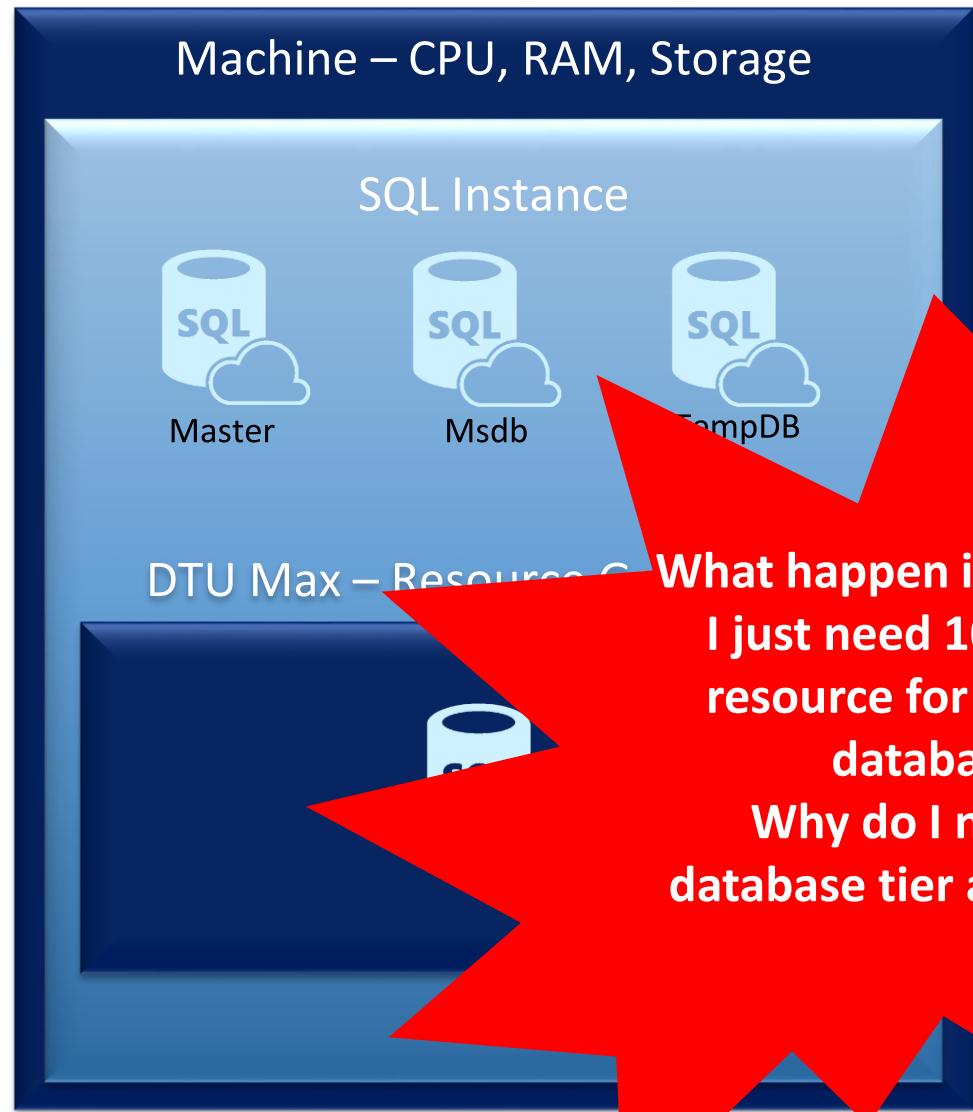
- ✓ Basic/Standard vs Web/Business
- ✓ Premium P6 y P11
- ✓ Elastic Database Pool

	Basic	Standard				Premium				
		S0	S1	S2	S3	P1	P2	P4	P6/P3	P11
Maximum database size	2 GB	250 GB			500 GB			1 TB		
DTUs	5	10	20	50	100	125	250	500	1,000	1,750
Point-in-time restore	Any point last 7 days	Any point last 14 days			Any point last 35 days					
Disaster recovery	Geo-Restore, restore to any Azure region	Standard Geo-Replication, offline secondary			Active Geo-Replication, up to 4 online (readable) secondary backups					
Max In-Memory OLTP storage	NA	NA	NA	NA	NA	1 GB	2 GB	3 GB*	8 GB	10 GB*
Max concurrent requests	30	60	90	120	200	200	400	800	1,600	2,400
Max concurrent logins	30	60	90	120	200	200	400	800	1,600	2,400
Max sessions	300	600	900	1,200	2,400	2,400	4,800	9,600	19,200	32,000

* In-Memory OLTP storage limits will soon adjust to 4 for P4 and 14 for P11.



SQL Azure Database V12 - Single Database



**What happen if all the time
I just need 10%-20% of
resource for a multiple
databases?
Why do I need this
database tier all the time?**

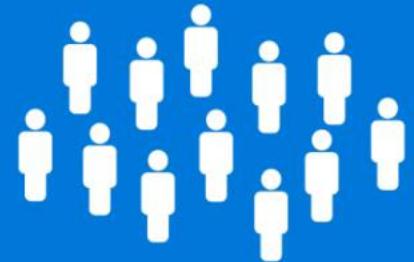
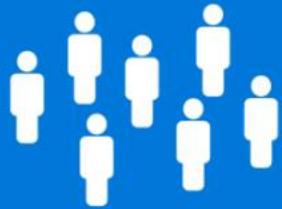
Single Tier has a price and resources

Application workload = should be

Single Tier

Can be incremented depending on #

databases.



1

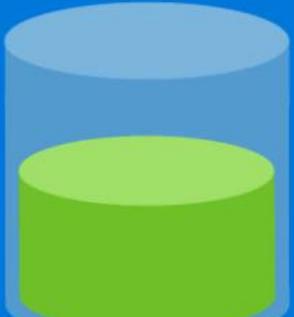
1

1

1



S0



S1



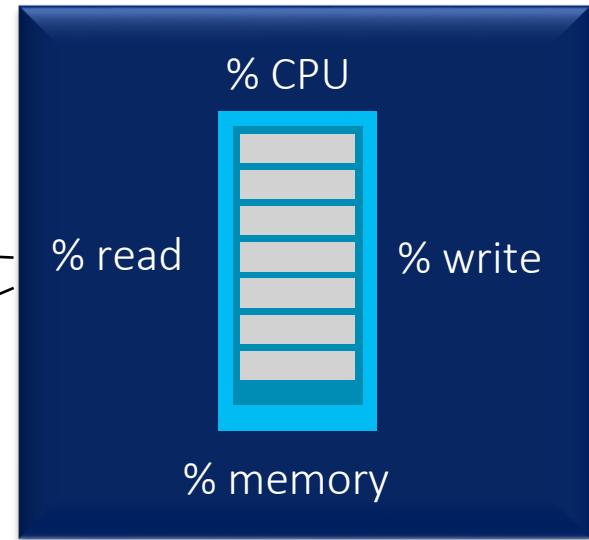
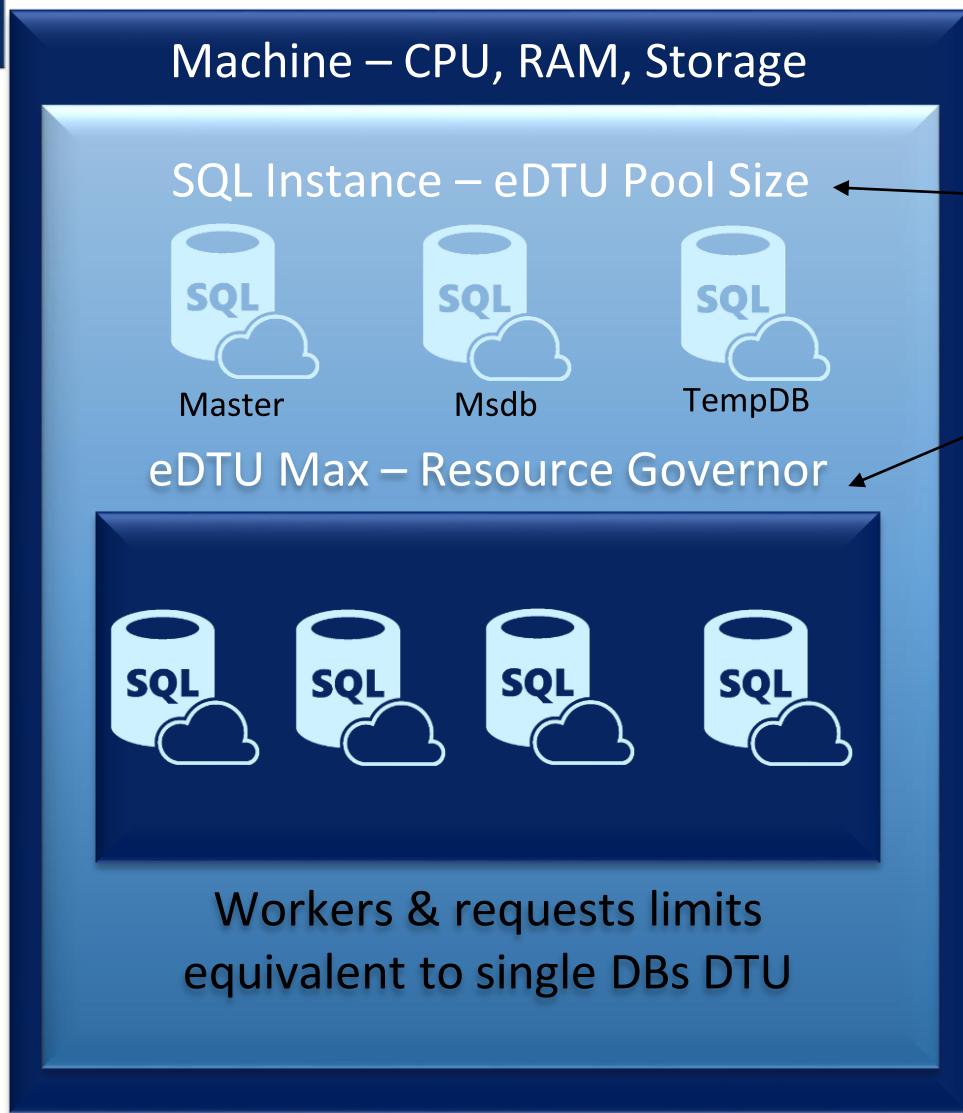
S2



S3



SQL Azure Database V12 – Elastic DB



- Every Database is able to use the eDTU DB Max.
- Good for customer with unpredictable performance or SaaS.
- The cost will be incremented depending on # databases and Pool Size

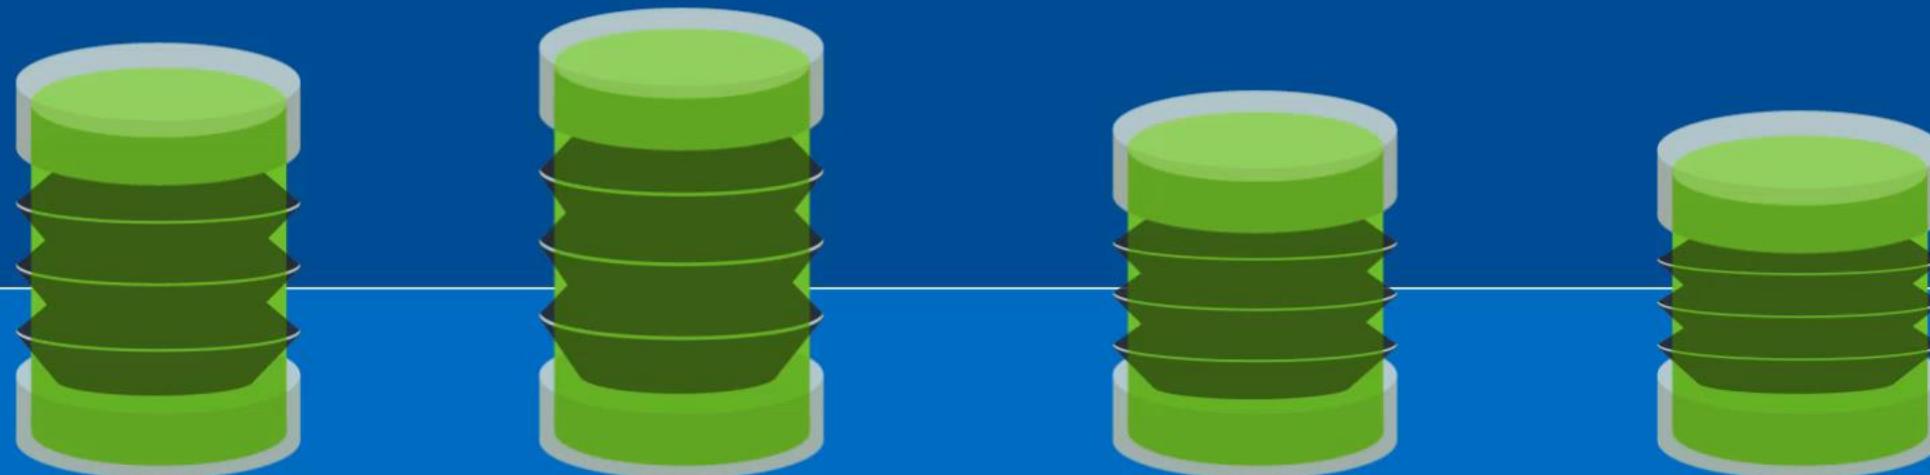
ELASTIC DATABASE POOL

MAX

100 DTUs

MIN

10 DTUs





SQL Azure Database V12 – Elastic Database

- A pool is given a set number of eDTUs, for a set price.

eDTU consumption

Individual databases are given the flexibility to auto-scale within set parameters.

Under heavy load a database can consume more eDTUs to meet demand.

Databases under light loads consume less, and databases under no load don't consume any eDTUs.



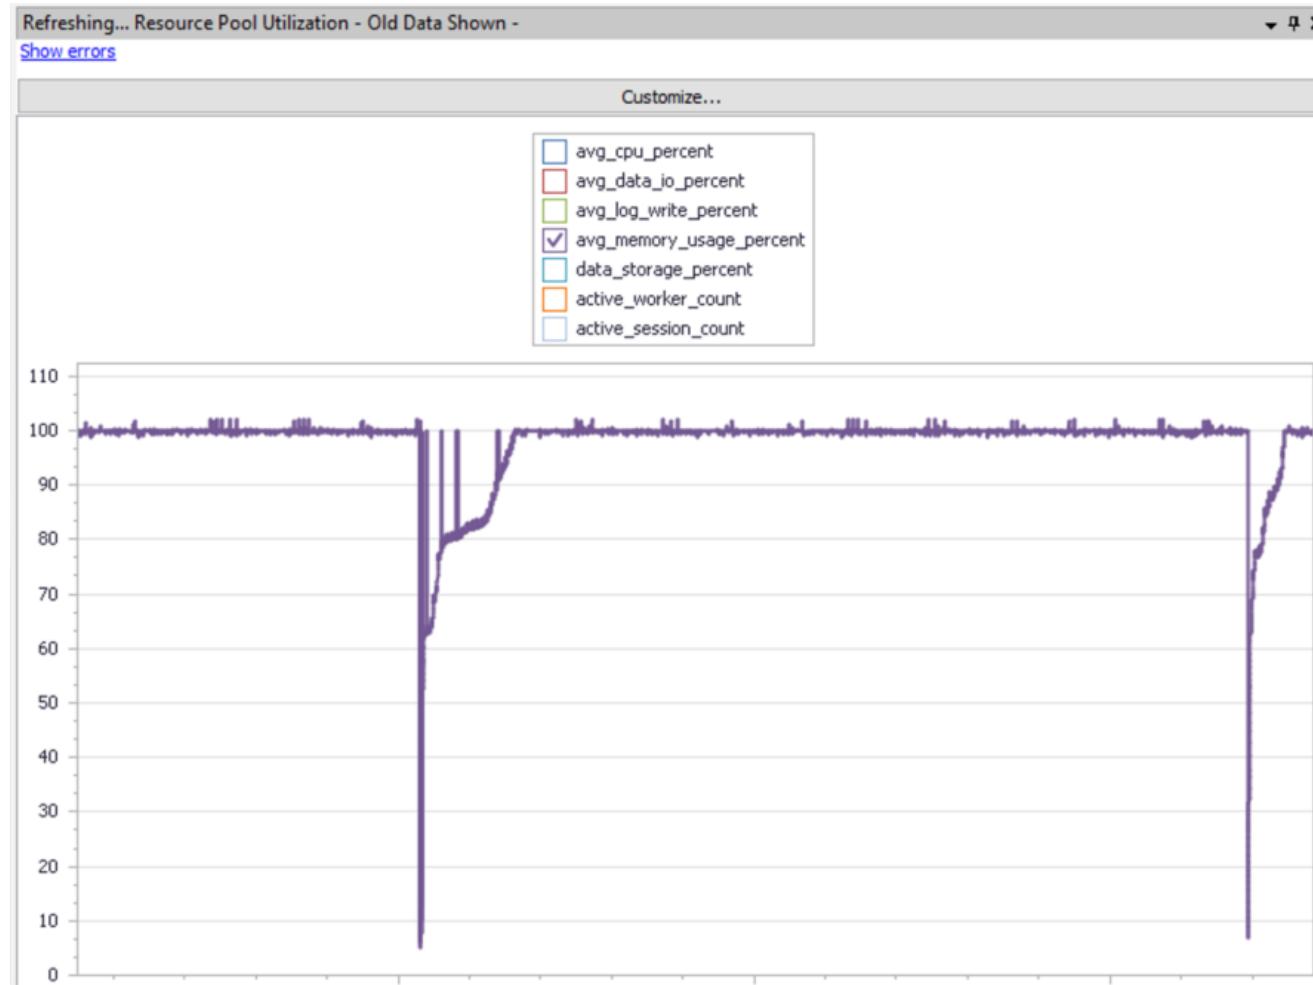
Provisioning resources for the entire pool rather than for single databases simplifies your management tasks. Plus you have a predictable budget for the pool.



SQL Azure Database V12 – Elastic Database - Think that...

Reconfiguration/failover impact

- Slow memory cached Warm-Up for 63 DBs running 200 eDTU/100 eDTU max





SQL Azure Database V12 – Elastic Database – Choosing Db Tier

Analysis From VM or On-premise

PSSDIAG for SQL Server
SQL Nexus.
Performance counters.
Reviewed instance configuration.
Azure SQL database calculator tool.
<http://dtucalculator.azurewebsites.net/>

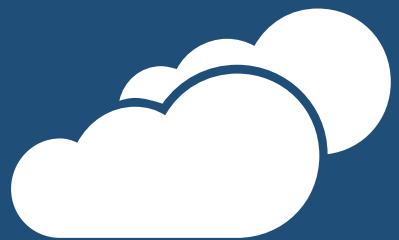
Analysis from Single SQL DBs

Elastic pool creation experience.





SQL Azure Database V12 – Elastic Database



Demo: Using Dtu-Calculator





SQL Azure Database V12 – Elastic Database

Elastic Tools



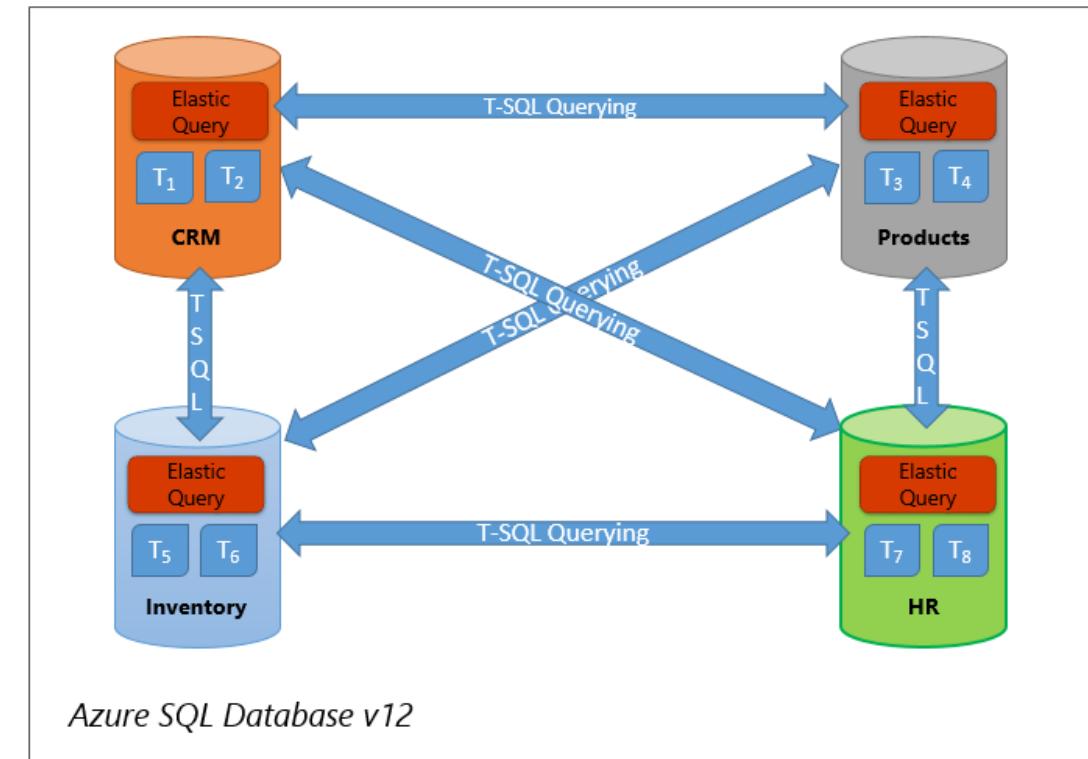
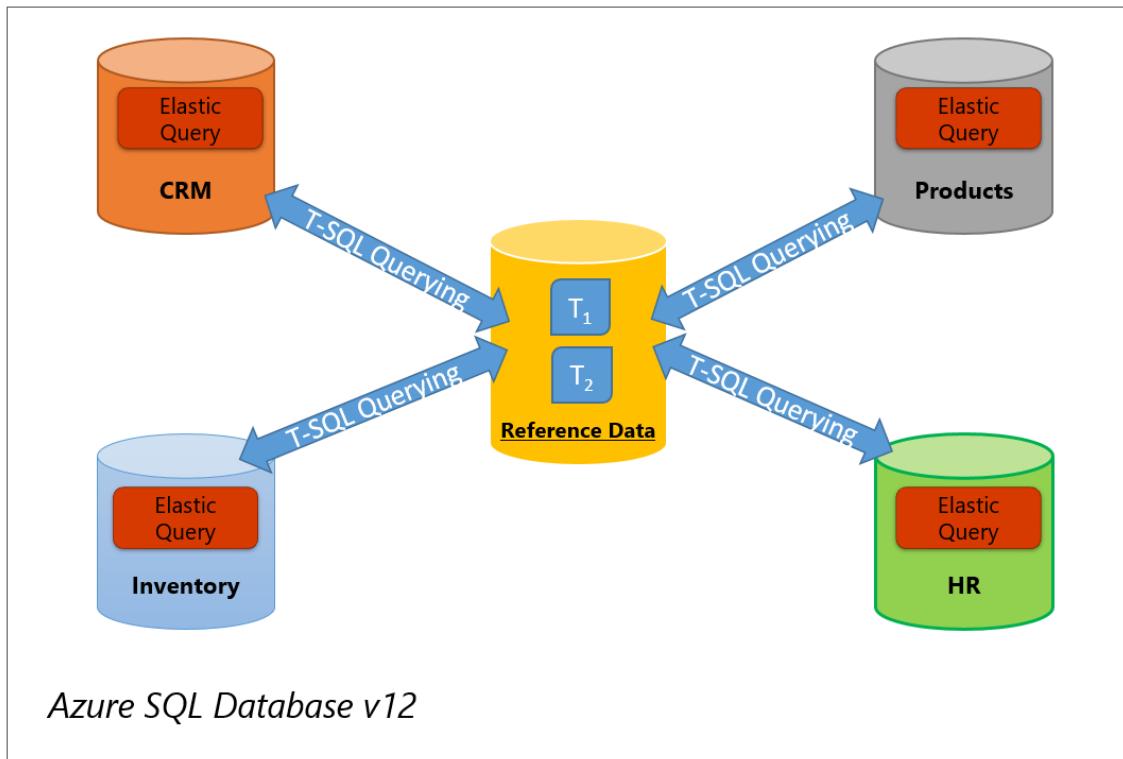
SQL Azure Database V12 – Elastic database query

- ✓ **Cross-database** query: read-only querying of remote databases.
- ✓ **Migrate** applications: using three- and four-part names or linked server to SQL DB.
- ✓ **Push SQL parameters**: to remote databases
- ✓ **Sp_execute_fanout**: parameters similar to sp_executesql
 - [Announcement blog post](#):
 - [Step-by-step tutorial](#)
 - [Overview documentation](#)



SQL Azure Database V12 – Elastic database query

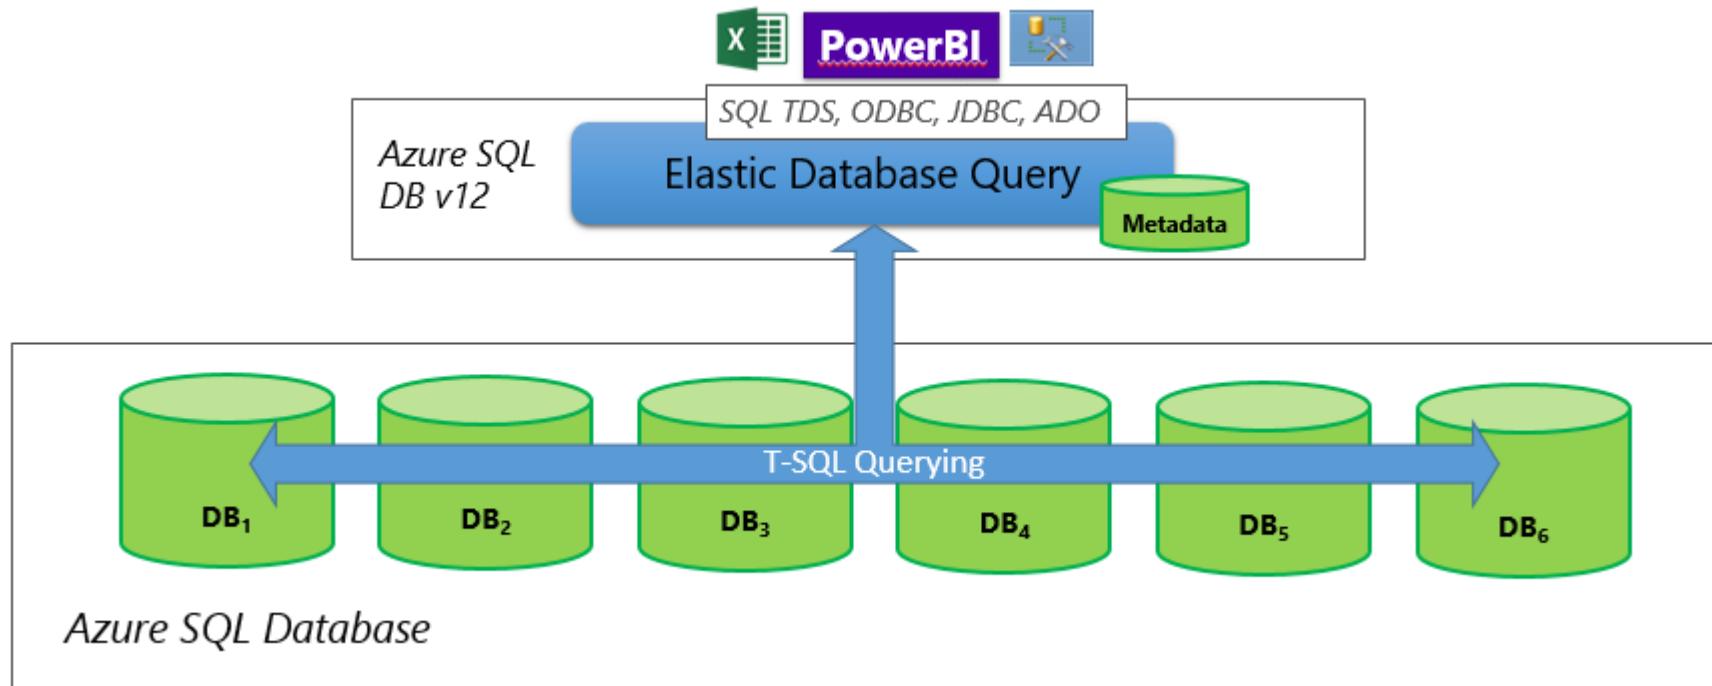
✓ Vertical partitioning





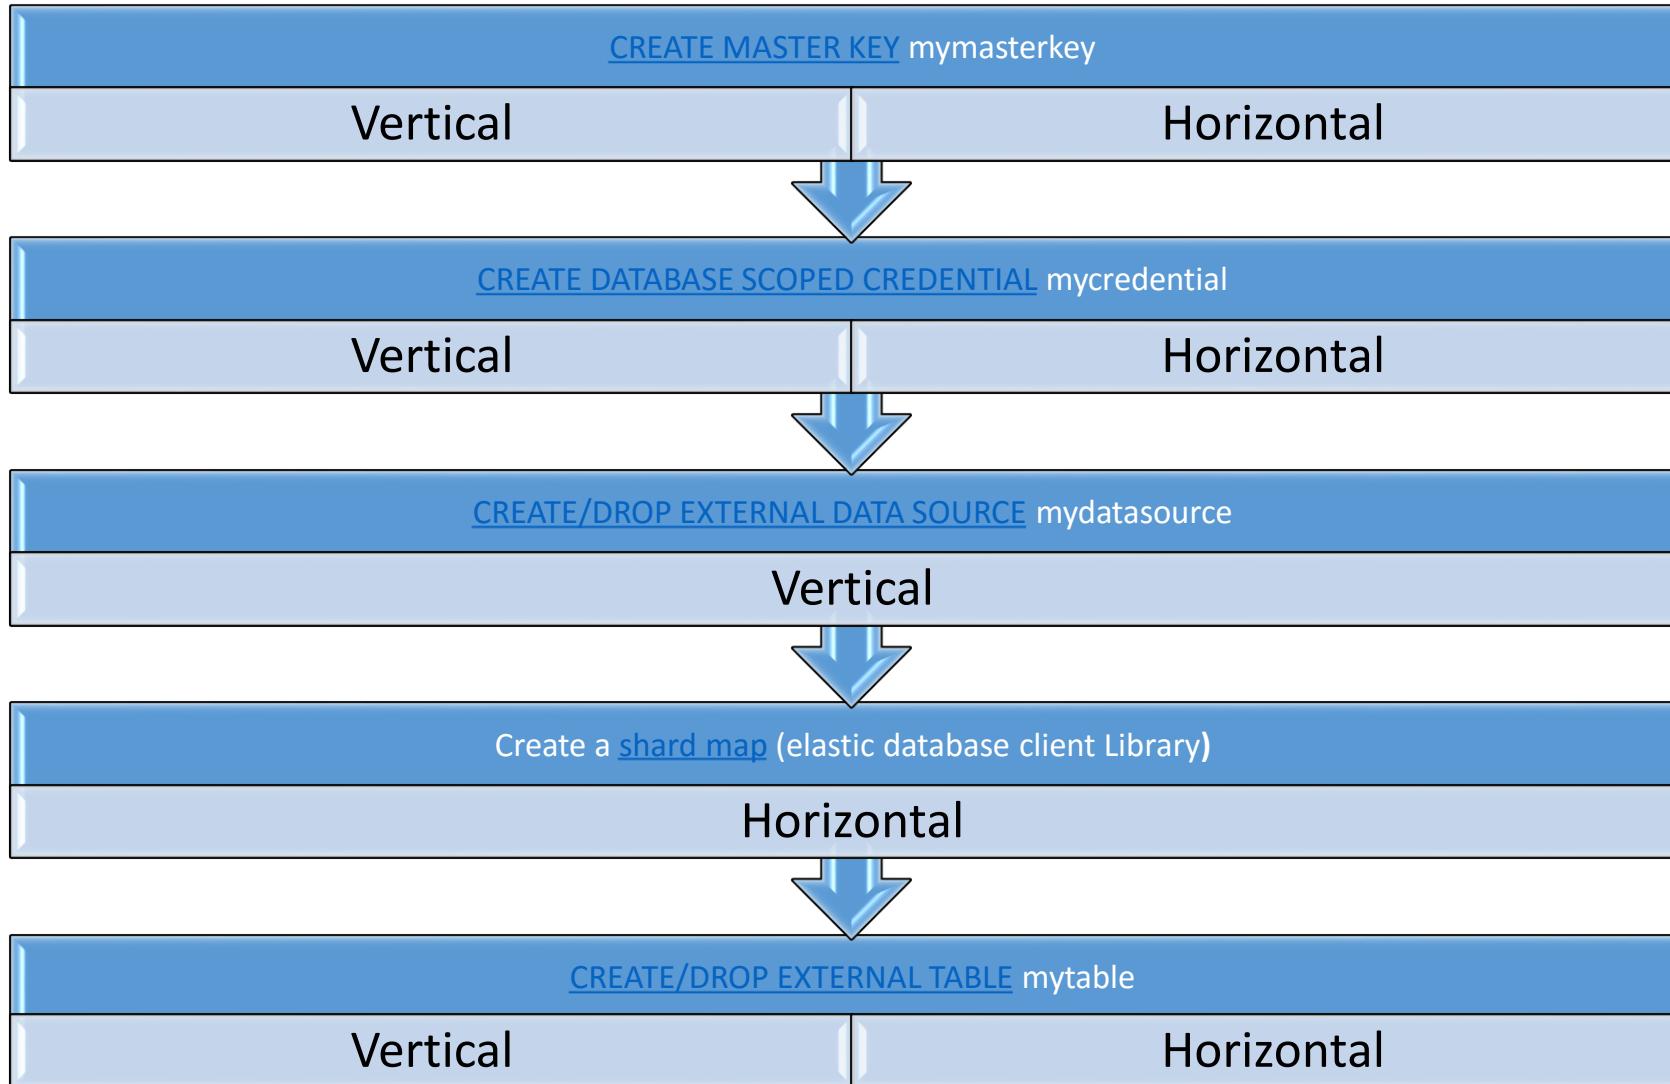
SQL Azure Database V12 – Elastic database query

✓ Horizontal partitioning



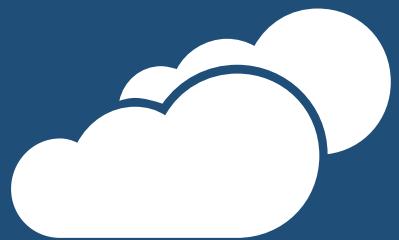


SQL Azure Database V12 – Elastic database query





SQL Azure Database V12 – Elastic database query



Demo: Elastic Query





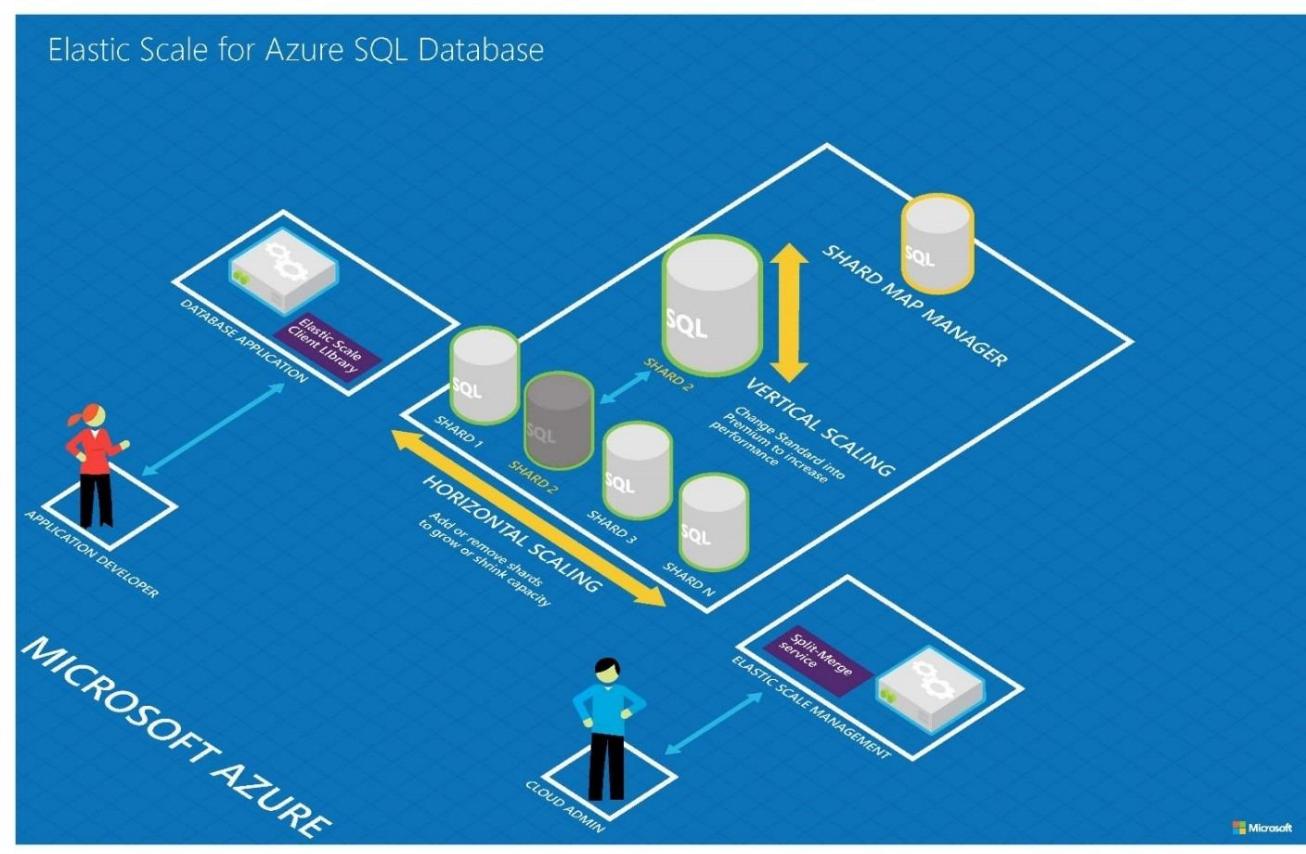
SQL Azure Database V12 - Elastic Scale

- ✓ The new Elastic Scale capabilities simplify the process of scaling out (and in) a cloud application's data tier.
- ✓ Elastic Scale offers functionality to help you build and maintain applications that can span a few databases or thousands.
- ✓ **High-volume OLTP:** Single application process massive data volumes by employing #number of database shards.
- ✓ **Multi-tenant SaaS:** Each customer is assigned their own database for higher isolation
- ✓ **Continuous data collection:** New shards for new date ranges. Newer shards can use higher service tiers, and scale down over time as usage is reduced





SQL Azure Database V12 - Elastic Scale





SQL Azure Database V12 - Elastic jobs

- ✓ Elastic Database jobs is a customer-hosted Azure Cloud Service that enables the execution of ad-hoc and scheduled administrative tasks, which are called jobs.
- ✓ **Administrative:** Deploy new schema
- ✓ **Update** data: product information common across all databases
- ✓ **Schedule:** automate updates after hours
- ✓ **Rebuild** indexes: regularly across a collection of databases
- ✓ **Centralize** query results: from many databases into a central table
- ✓ **Reference:** <https://azure.microsoft.com/en-us/documentation/articles/sql-database-elastic-jobs-powershell/>
- ✓ **Service:** <https://azure.microsoft.com/en-us/documentation/articles/sql-database-elastic-jobs-service-installation/>
- ✓ **Nuget:** <https://www.nuget.org/packages/Microsoft.Azure.SqlDatabase.Jobs/>



Elastic jobs



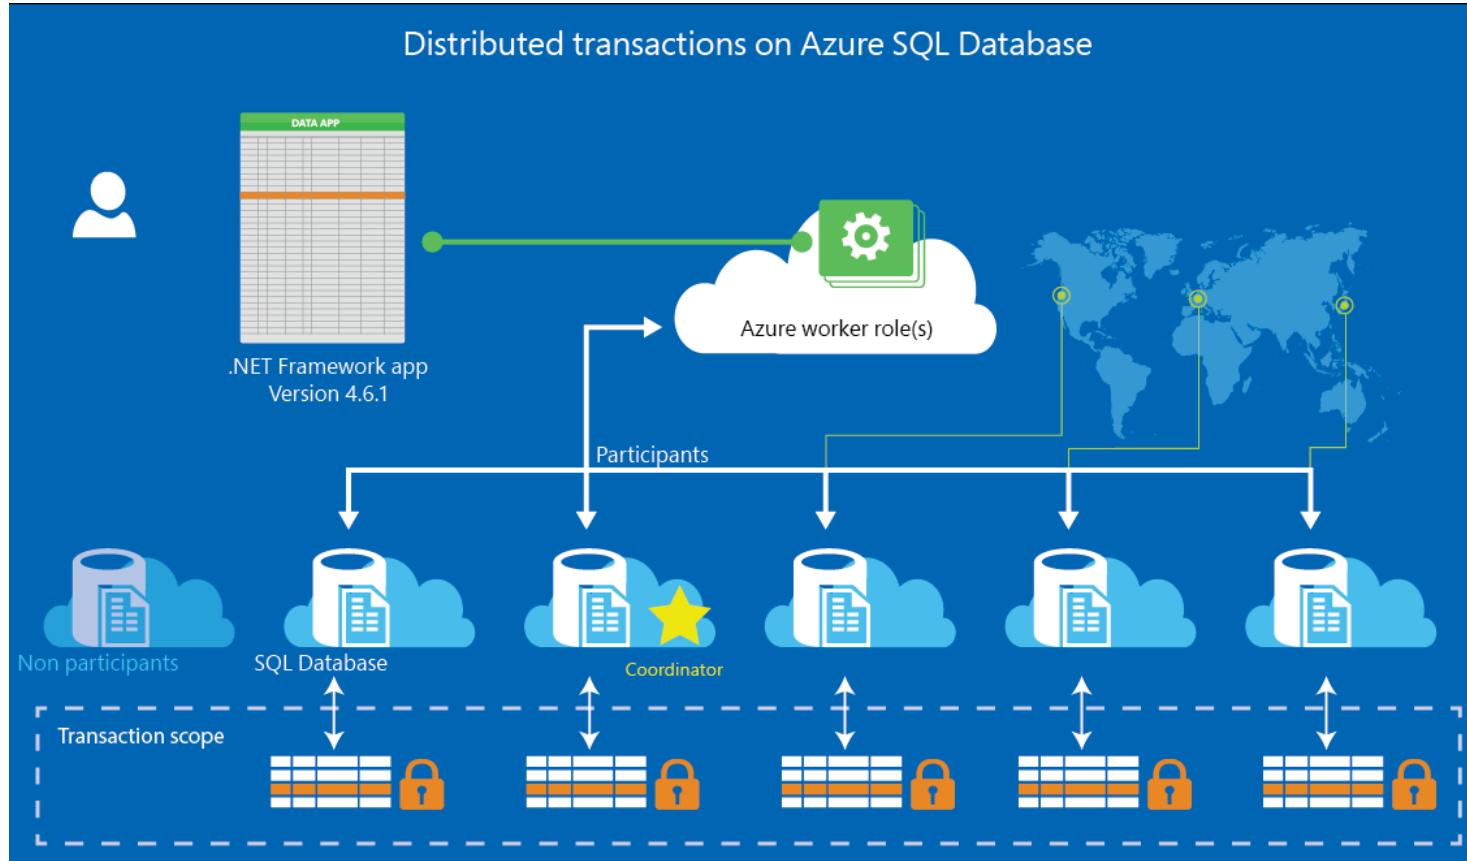


SQL Azure Database V12 - Elastic Database Transactions

- ✓ Elastic database transactions for Azure SQL Database (SQL DB) allow you to run transactions that span several databases in SQL DB
- ✓ Available for **.NET 4.6.1** applications: ADO .NET ->System.Transaction classes
- ✓ Coordinate distributed transactions is **directly integrated** into SQL DB As MSDTC is not an Azure PaaS.
- ✓ Applications: Connect to any SQL Database-> Launch distributed transactions-> one of the databases will coordinate **transparently**
- ✓ **Multi-database** applications: Vertically partition
- ✓ **Sharded database** applications: Horizontal partitioning
- ✓ **Only SQL DB supported**
 - ✓ Not Other [X/Open XA](#) resource providers
 - ✓ Not SQL Server
- ✓ **Only client-coordinated** transactions
- ✓ Only databases on **Azure SQL DB V12** are supported.



SQL Azure Database V12 - Elastic Database Transactions





SQL Azure Database V12 - Elastic Database Transactions

```
using (var scope = new TransactionScope())      {  
    using (var conn1 = new SqlConnection(connStrDb1))  
    {  
        conn1.Open();  
        SqlCommand cmd1 = conn1.CreateCommand();  
        cmd1.CommandText = string.Format("insert into T1 values(1)");  
        cmd1.ExecuteNonQuery();  
    }  
    using (var conn2 = new SqlConnection(connStrDb2))      {  
        conn2.Open();  
        var cmd2 = conn2.CreateCommand();  
        cmd2.CommandText = string.Format("insert into T2 values(2)");  
        cmd2.ExecuteNonQuery();  
    }  
    scope.Complete();
```





SQL Azure Database V12 - Elastic Database Transactions



Demo: Elastic Transaction





SQL Azure Database V12

Connectivity



What is new using V12 Azure SQL Connectivity

.Net Framework

[.NET 4.5.1 or above:](#)

[ConnectRetryCount and](#)

[ConnectRetryInterval](#)

[.NET 4.6.1, retry logic implemented at SqlConnection.Open](#)

[Entity Framework 6.0:](#)

[SqlAzureExecutionStrategy](#)

To measure

[SqlConnection object properties](#)

[Stopwatch](#)

[EF: Optimizar query performance and modelo](#)

Tips & Tricks

Documentation on handling connections to Azure SQL Db:

<https://azure.microsoft.com/en-us/documentation/articles/sql-database-connect-central-recommendations>

Retry logic for Azure SQL Db:

<http://social.technet.microsoft.com/wiki/contents/articles/4235.retry-logic-for-transient-failures-in-windows-azure-sql-database.aspx>

Error Code List for Retry Logic in Azure SQL Db: <https://azure.microsoft.com/en-us/documentation/articles/sql-database-develop-error-messages/>

`SqlConnection.ClearPool(cnn)` - To clean up the connection from the connection pool
`using SqlConnection.ClearPool(cnn)` -

<http://blogs.msdn.com/b/bartr/archive/2010/06/18/sql-azure-connection-retry.aspx>

New connectivity method

Direct connection, not proxy

Differences depending on data provider.





SQL Azure Database V12 - StopWatch and SqlConnection

```
//Call StopWatch function  
Stopwatch stopWatch = new Stopwatch();  
    stopWatch.Start();  
    stopWatch.Stop();  
//Elapsed time TimeSpan.  
    TimeSpan ts = stopWatch.Elapsed;  
//Obtained the time spent and data from SqlConnection property.  
IDictionary currentStatistics = awConnection.RetrieveStatistics();  
long bytesReceived = (long)currentStatistics["BytesReceived"];  
    long bytesSent = (long)currentStatistics["BytesSent"];  
    long selectCount = (long)currentStatistics["SelectCount"];  
    long selectRows = (long)currentStatistics["SelectRows"];
```



SQL Azure Database V12 – Connectivity



Demo: SQL Connection Properties





SQL Azure Database V12 – Retry Policy

- ✓ Caches physical connections with the same conn string in the process
- ✓ Reduces the cost of opening/closing operations
- ✓ If a connection fails to connect, it will be dropped from the pool
- ✓ What if a connection succeed to connect, but fails to execute commands?
 - ✓ It may remain in the pool for a long time and create additional issues
- ✓ Workaround: after multiple failed retries (e.g. 3) on the same command execution, clean up that connection from the pool
 - ✓ `SqlConnection.ClearPool(cnn)` - **To clean up the connection from the connection pool using `SqlConnection.ClearPool(cnn)`** - <http://blogs.msdn.com/b/bartr/archive/2010/06/18/sql-azure-connection-retry.aspx>



SQL Azure Database V12 – Retry Policy

- ✓ Open connection late and close asap (especially disposable resources)
- ✓ Leverage connection pooling
- ✓ Retrieve only the data you need
- ✓ Cache metadata and data where possible, avoid metadata retrieval at runtime
- ✓ If possible, batch SQL statements together
- ✓ Use Stored Procedures especially to encapsulate complex data access logic and transactional behaviors, this will naturally reduce round trips and resource locking on server side
- ✓ Leverage Table Value Parameters
- ✓ Avoid application based transaction management
- ✓ Keep transactions as short as possible, avoid to involve unnecessary resources, use appropriate isolation level



SQL Azure Database V12 – Retry Policy

```
int retryIntervalSeconds = 10;  
bool returnBool = false;  
  
for (int tries = 1; tries <= 5; tries++)  
{  
    try  
    {  
        if (tries > 1)  
        {  
            H.Thread.Sleep(1000 * retryIntervalSeconds);  
            retryIntervalSeconds = Convert.ToInt32(retryIntervalSeconds * 1.5);  
            C.SqlConnection.ClearAllPools();  
        }  
        sqlConnection.ConnectionString = clsConexion;  
        sqlConnection.Open();....  
  
        catch (C.SqlException sqlExc)  
        {  
            if (sqlExc.Number == 4060 || sqlExc.Number == 40197....
```





SQL Azure Database V12

New Tools



SQL Azure Database V12 - Common Performance Issues

- **Performance issues**

- ✓ Query tuning, investigating execution plans.
- ✓ Query Timeouts.
- ✓ Blocking.
- ✓ Deadlocks?.
- ✓ Database Tier in Azure SQL Database

- **Frequently Simple**

- ✓ Outdated statistics.
- ✓ Missing Indexes.
- ✓ Throttling for Azure SQL Database.

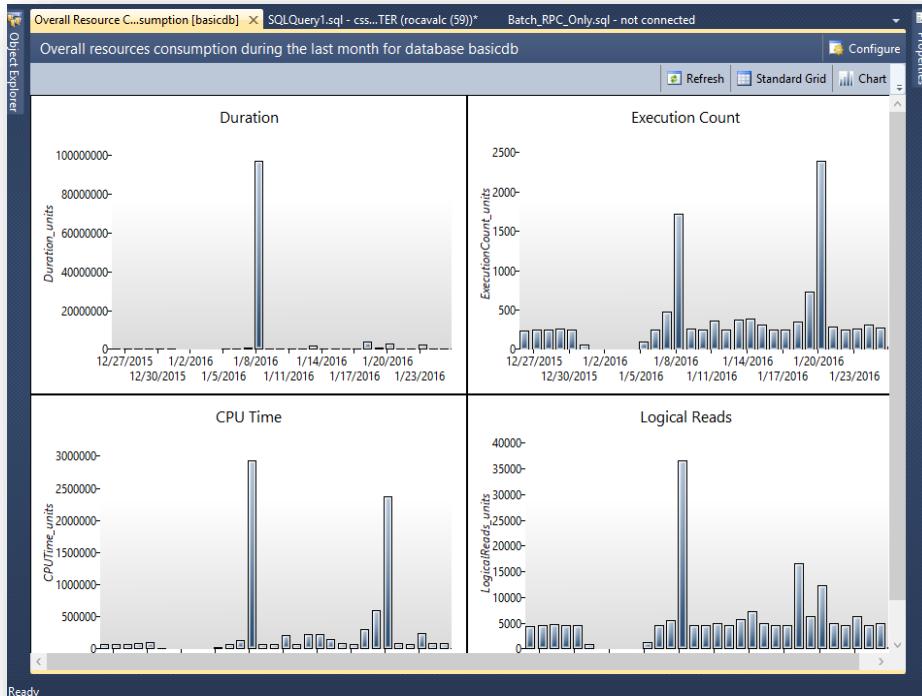
Way to resolve

- ✓ Blocking - SDP, SQL Nexus
- ✓ Attentions, Errors, etc. – Xevents.
- ✓ Query behavior – QDS, DMVs

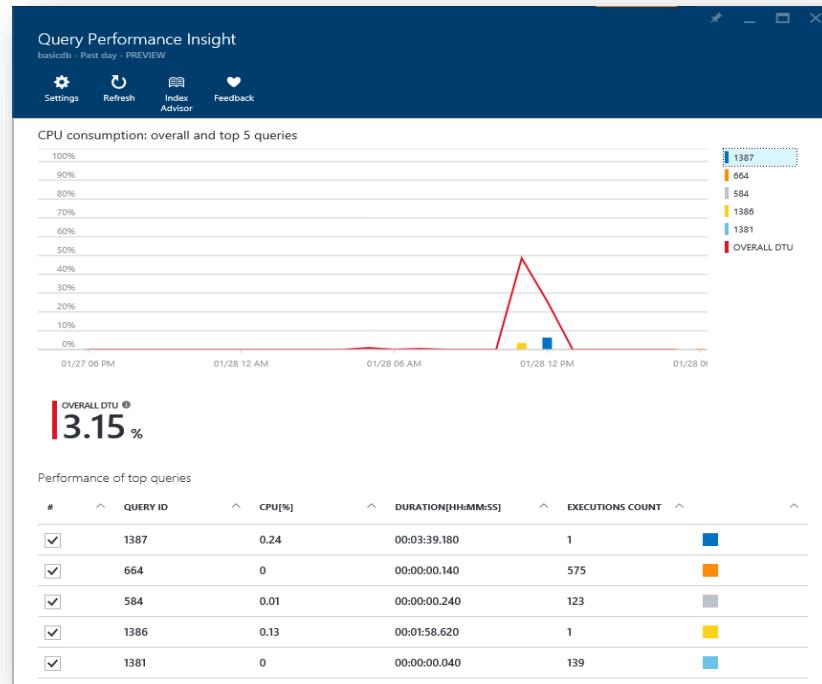




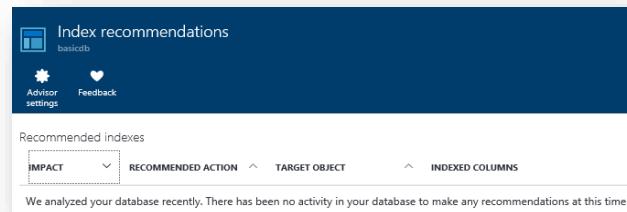
What is new using V12 Azure SQL Tools



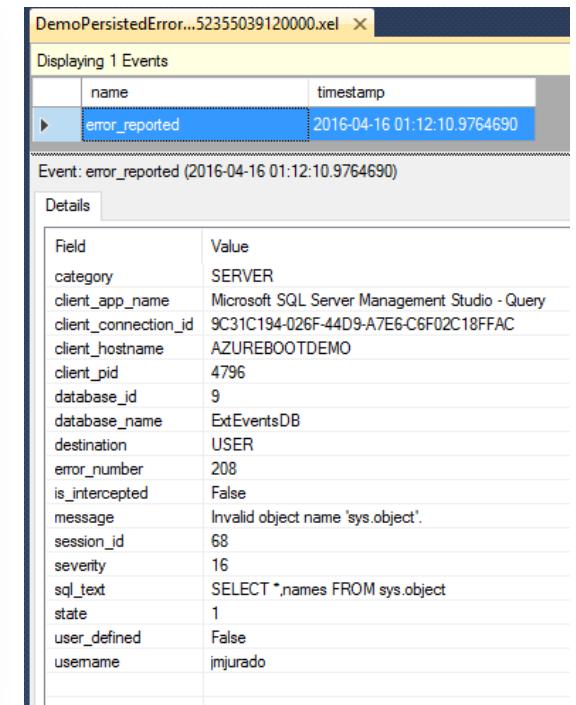
Query Data Store



Query Performance Insights



Index Advisor



Extended Events





SQL Azure Database V12 - Query Data Store

- ✓ **It is a *persisted database with query execution information for***
 - ✓ SQL Server 2016 (Preview)
 - ✓ Azure SQL DB V12
- ✓ **Query Performance Tuning and Troubleshooting**
 - ✓ SQL Profiler replacement in some parts.
 - ✓ Minimum impact for SQL Engine in SQL Server or Azure SQL DB.
 - ✓ Other tools like Query Performance Insight, Index Advisor are using QDS Information.
 - ✓ SSMS Reports Supported for SQL Server and Azure SQL Database.
 - ✓ Catalog Views all available.
 - ✓ Active or disable this feature demand.
 - ✓ Use Extended Events to capture the information.





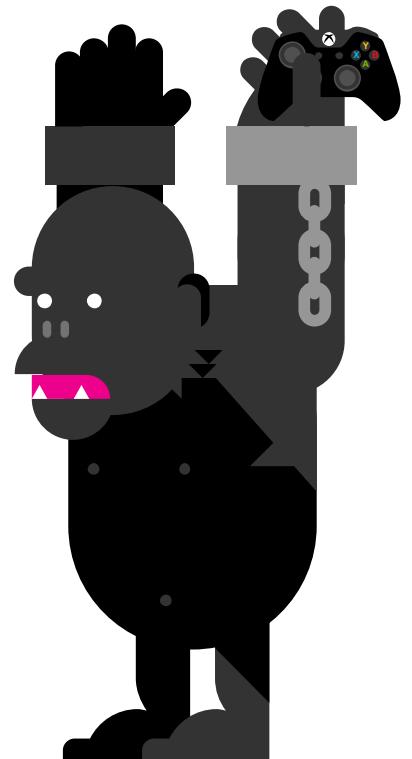
SQL Azure Database V12 - Information collected

✓ **A Query Flight Recorder to store:**

- ✓ Queries/statements, Query Plans, Runtime statistics with connection properties.
- ✓ Maintains the information even an instance restart.
- ✓ Just save DML: SELECT, DELETE, INSERT, UPDATE.
- ✓ Doesn't record details about: DDL: BULK INSERT or DBCC commands.
- ✓ The information will be recorded with a limited time.
- ✓ Details of queries whether they are cached or not.
- ✓ Analyze the history
- ✓ Look at properties and statistics for queries not available in DMVs
- ✓ Allows a simpler and more robust method to force plans vs plan guides
- ✓ Independent of what is cached
- ✓ Statement text appears in parameterized form
- ✓ Including cursors and SET assignment
- ✓ Persistence is async
- ✓ Works with In-Memory tables in Azure SQL DB and SQL SERVER 2016.

✓ **Other Information**

- ✓ Use the scenarios [page](#)
- ✓ The docs. Start [here](#)
- ✓ Query Store [Best Practices](#)

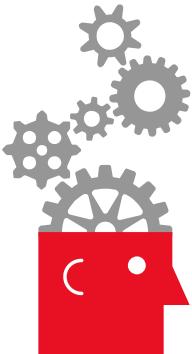




SQL Azure Database V12 - CAPTURE MODE

- **Query Capture Mode**

- ✓ Query Store supports adjustable query capture.
 - ✓ The goal of capture option is to provide users with additional control over ingress of new queries with main value prop to enable collection of data that is highly relevant.
- ✓ This is the list of supported values:
 - ✓ ALL - all queries are captured. Infrequent queries and queries with insignificant compile and execution duration are ignored. Thresholds for execution count, compile and runtime duration are internally determined.
 - ✓ AUTO - capture relevant queries based on execution count and resource consumption. This is the default configuration value
 - ✓ NONE: stop capturing new queries. Query Store will continue to collect compile and runtime statistics for queries that were captured already.
- ✓ To change capture mode use one of the following statements:
 - ✓ ALTER DATABASE <database_name> SET QUERY_STORE (QUERY_CAPTURE_MODE = AUTO);
 - ✓ **Query capture mode** of AUTO (Default is ALL) = capture relevant queries based on execution count and resource consumption.
NONE = Stop capture new queries but track and persist stats for existing.
 - ✓ ALTER DATABASE <database_name> SET QUERY_STORE (QUERY_CAPTURE_MODE = NONE);
 - ✓ **Sized based cleanup** of AUTO (Default is OFF) = Remove oldest queries and least expensive when 90% of max (until 80% reached).
Overrides time retention (days kept in store)
- ✓ Query capture = AUTO; Query retention = 30 days; Sized based cleanup = ON
- ✓ We do limit memory for Query Store and may revert to read-only mode memory overcommitted or when the disk is full.





SQL Azure Database V12 - QDS vs DMVs

✓ Using sys.query_context_settings

- ✓ set_options
- ✓ cursor_options - Go [here](#) to interpret bitmask
- ✓ schema_id

✓ Comparing DMVs vs Query Store

- ✓ Compilation statistics
- ✓ Plan properties including engine version and database compatibility level
- ✓ Details and stats for plan variations used for a query over time
- ✓ Recorded as queries are compiled
- ✓ Record query execution stats ***even if aborted***





Dynamic Management Views

Plan store Runtime stats

sys.database_query_store_options

- Exposes various Query Store configuration options for user database.

sys.query_store_query_text

- Exposes information about captured query text for a query

sys.query_context_settings

- Exposes information about content settings

sys.query_store_query

- Exposes information about queries in every execution

sys.query_store_plan

- Exposes information about different plans SQL Server used to execute queries in the system.

sys.query_store_runtime_stats_interval

- Returns all intervals created in Query Store

sys.query_store_runtime_stats

- Returns runtime statistics for executed query plans, aggregated on per-interval basis.





Dynamic Management Views

Force plan

- `sp_query_store_force_plan @query_id, @plan_id`

Stop plan forcing

- `sp_query_store_unforce_plan @query_id, @plan_id`

Clear all data for a single query

- `sp_query_store_remove_query @query_id`

Remove a single plan

- `sp_query_store_remove_plan @plan_id`

Clear all stats

- `sp_query_store_reset_exec_stats @plan_id`

Flush data

- `sp_query_store_flush_db`





SQL Azure Database V12 - Maintenance

- Enabling, Clearing, and State

- ✓ **ON** = Enable; **OFF** = Disabled (existing state and data kept); **CLEAR** = TRUNCATE tables
- ✓ **READ_WRITE** = Default when ON; **READ_ONLY** = intentional or *problem* (desired != actual)

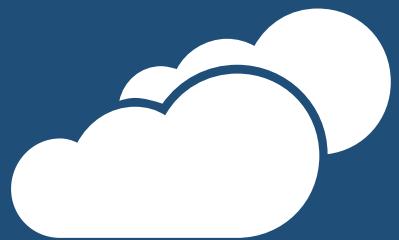
- Size, limits, and retention

- ✓ Default **max size** = 100Mb (limited by overall database size). If you hit max state = READ_ONLY
- ✓ Default **max plans per query** = 200
- ✓ Default **days queries kept in store** = 367 days





Azure SQL Database – Create First DB



Demo: Elastic Query





SQL Azure Database V12 - Extended Events

- You have the documentation in this [URL](#)
- In order to use Extended Events with your SQL Database server there are three basic steps:
 - Setup a storage account and container
 - Enable access to storage from SQL Database
 - Collect Data using Extended Events

Setup storage account and container

XEvents use blob storage as their target. You shall need access to a storage account that will be the target for writing out the XEL files. Typically, this shows up under the Storage node in Management Portal. For example:

A screenshot of the Azure Storage blade. It shows a list of storage accounts with one item: 'cssxestore'. The account is marked as 'Online' and is located in 'South Central US'. There is a 'Manage' button next to it.

A screenshot of the 'Manage Access Keys' dialog box. It contains fields for 'STORAGE ACCOUNT NAME' (set to 'cssxestore'), 'PRIMARY ACCESS KEY' (a long yellow-highlighted string), and 'SECONDARY ACCESS KEY' (another long string). Each key has a 'regenerate' button to its right. A note at the top says: 'When you regenerate your storage access keys, you need to update any virtual machines, media services, or applications that access this storage account to use the new keys.' A 'Learn more' link is provided.

After storage is created, note the "Storage Account Name" and the "Primary Access Key". Both will be needed to access the .xel output files within the storage.



SQL Azure Database V12 - Extended Events

- Enable access to storage from SQL Database.
- You shall use a tool called Azure Storage Explorer to download the XEL files from the container to your local machine. Download and install Azure Storage Explorer (<http://azurestorageexplorer.codeplex.com/>) on your client machine.
- We used the current version (Version 6) as of this writing; feel free to use the latest version if another one is available.
 - Here are the steps to do for the same:
 - Configure Azure Storage Explorer to point to your storage account using the "Storage Account Name" and "Storage Account Key" from previous section titled Setup storage account and container. Check the Use HTTPS" checkbox and then click on Test Access





SQL Azure Database V12 - Extended Events

- The “Test Access” button click should return a success message. Click Save to save the configuration



You should see a similar screen to this once you are done



The next step is to add a container within the storage account. Click Blob Containers and add a new container



In my case I created a container named xe_container (with access level=off)





SQL Azure Database V12 - Extended Events

In my case I created a container named xe_container (with access level=off)



that show up like this



Next click on the specific container and then click on Security



Go to Shared Access Signatures tab and choose Read and Write Permission and ensure that End Date is well in future so the access does not expire before data collection is done.





SQL Azure Database V12 - Extended Events

Go to Shared Access Signatures tab and choose Read and Write Permission and ensure that End Date is well in future so the access does not expire before data collection is done.



Select the "Generate Signature" button, then copy the output to notepad



The output of this step should be a SAS key for the container that looks like this:

http://your_storage_account.blob.core.windows.net/your_container_name?sr=c&sv=start_date&sr=c&sig=%signature Microsoft



SQL Azure Database V12 - Extended Events

- Collecting Data using Extended Events

Open the SQL SERVER Management Studio and follow up the following instructions:

- Create a master key specifying a strong password

```
CREATE MASTER KEY ENCRYPTION BY PASSWORD = 'xxxxxx!';
```

- Define the blob storage where the XEL will be save. Use the SAS key that was provided by the definition of the blob storage.

```
CREATE DATABASE SCOPED CREDENTIAL [https://xxxx.blob.core.windows.net/xe-container]
```

```
WITH IDENTITY='SHARED ACCESS SIGNATURE',
```

```
SECRET = 'sv=2014-02-14&sr=c&sig=Hz2n9vs%3D&st=2016-01-25T23%3A00%3A00Z&se=2016-02-02T23%3A00%3A00Z&sp=rw'
```

- Define the Extended Event. The filter result <> 0 will capture when the error gives an error.

```
CREATE EVENT SESSION
```

```
ON DATABASE
```

```
ADD EVENT
```

```
    ACTION
```

```
    pid
```

```
WHERE
```

```
    ADD TARGET
```

```
        SET filename='https://xxx.blob.core.windows.net/xe-container/DemoPersistedEvents.xel' )
```

- Start the event and wait to reproduce the issue.

```
ALTER EVENT SESSION [ ] ON DATABASE STATE
```

- Once the issue has been reproduced, stop the event.

```
ALTER EVENT SESSION [ ] ON DATABASE STATE = STop;
```

- You could delete it when you don't need it anymore.

```
DROP EVENT SESSION [ ] ON DATABASE
```

- You should see XEL files in the storage container in Azure Storage Explorer

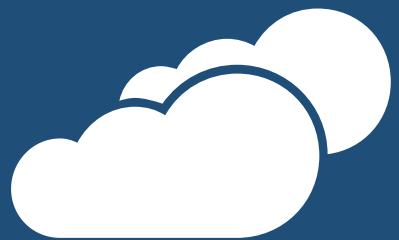
- You can then download to your laptop/local machine

- (Note that the files are in the format server_name_database_name_session_name.— so you can tell which server & database they were collected from)





Azure SQL Database – Create First DB



Demo: Elastic Query





SQL Azure Database V12 - DMVs

- ```
SELECT end_time , (SELECT Max(v) FROM (VALUES (avg_cpu_percent) , (avg_data_io_percent) , (avg_log_write_percent)) AS value(v)) AS [avg_DTU_percent] FROM sys.dm_db_resource_stats ORDER BY end_time DESC;
```
- Tip: To find a resource usage:
  - `sys.dm_db_resource_stats`. Has 15 second granularity.
  - `sys.resource_stats` has 5 minute granularity.
  - `sys.elastic_pool_resource_stats` for elastic pool databases.
  - `sys.dm_os_performance_counters` -  
<http://blogs.msdn.com/b/dfurman/archive/2015/04/02/collecting-performance-counter-values-in-sql-azure.aspx>



# SQL Azure Database V12 - Engine - Tips & Tricks

- Deadlock:

```
SELECT *, CAST(event_data as XML).value('(/event/@timestamp)[1]', 'datetime2') AS timestamp ,CAST(event_data as XML).value('(/event/data[@name="error"]/value)[1]', 'INT') AS error ,CAST(event_data as XML).value('(/event/data[@name="state"]/value)[1]', 'INT') AS state ,CAST(event_data as XML).value('(/event/data[@name="is_success"]/value)[1]', 'bit') AS is_success ,CAST(event_data as XML).value('(/event/data[@name="database_name"]/value)[1]', 'sysname') AS database_name FROM sys.fn_xe_telemetry_blob_target_read_file('el', null, null, null) where object_name = 'database_xml_deadlock_report'
```



# SQL Azure Database V12 - Code Best Practices – Tips & Tricks

- Batching considerations: <https://azure.microsoft.com/en-us/documentation/articles/sql-database-use-batching-to-improve-performance/>
- Performance considerations: <https://azure.microsoft.com/en-us/documentation/articles/sql-database-performance-guidance/>
- To avoid any parameter sniffing issue.  
<http://blogs.technet.com/b/mdegre/archive/2012/03/19/what-is-parameter-sniffing.aspx>



# SQL Azure Database V12 - Entity Framework - Tips & Tricks

- **Link describing Performance Considerations with Entity Framework – with regards to optimize query performance and model.** <http://msdn.microsoft.com/en-us/library/cc853327.aspx>



# SQL Azure Database V12

## *Alerts*



# SQL Azure Database V12 - Alerts and events

Define rule

Specify threshold

Threshold violated

Alert rule active

Registers an alert

Send E-mail (optional)



# SQL Azure Database V12 - Alert Rules

Blocked by Firewall

Failed Connections

Successful Connections

CPU Percentage

Deadlocks

DTU Percentage

Log IO Percentage

Data IO Percentage

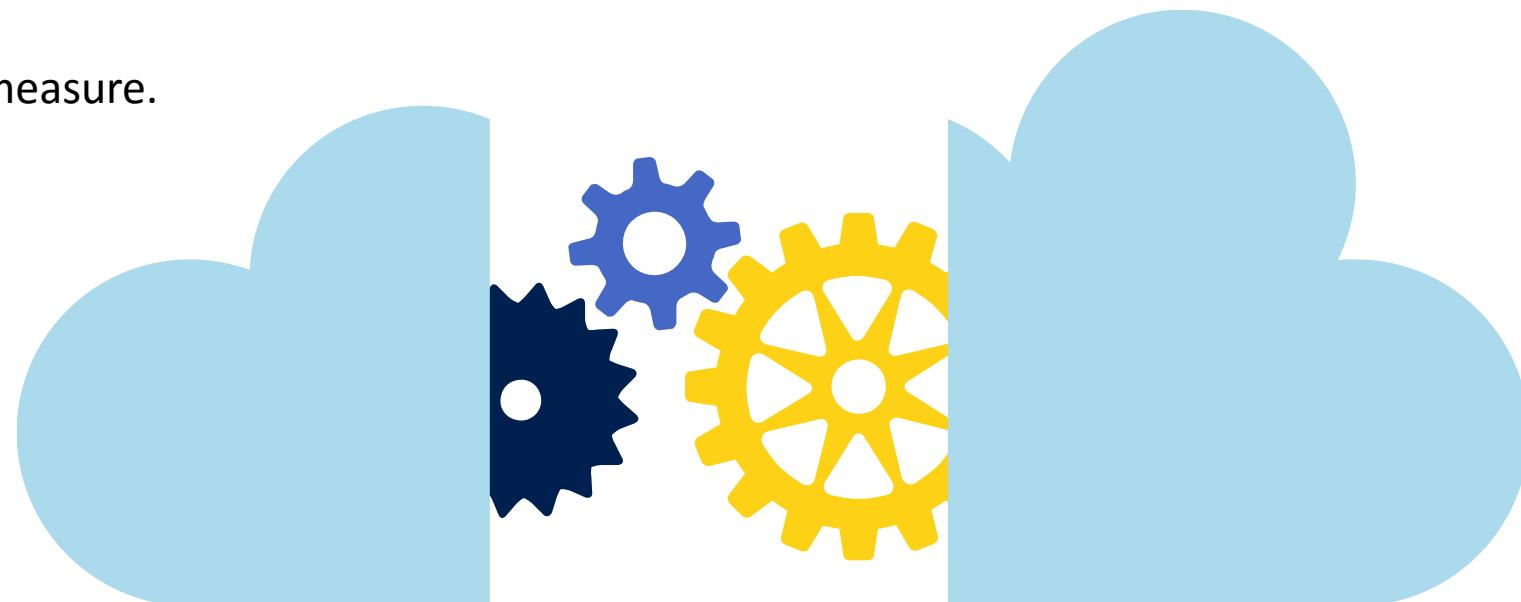
Total Database Size





# SQL Azure Database V12 - Alerts – 3rd Party tool

- New Relic issue causing a high memory usage, must be to review the actual tool that customer has.
- Review with Nagios some timeouts and our customer needs to implement other alternatives.
- Performance issue with you have a lot of row in the extended events.
- Review the unit of measure.





# SQL Azure Database V12

*Q&A*

If you have any question, please send an email

[jmjurado@microsoft.com](mailto:jmjurado@microsoft.com)





# SQL Azure Database V12

*Thanks!!*

